Evolution of application architecture

...

How did we get to service mesh?

Monolith application

Single unit of executable

=

Application

=

Single process

Application modules

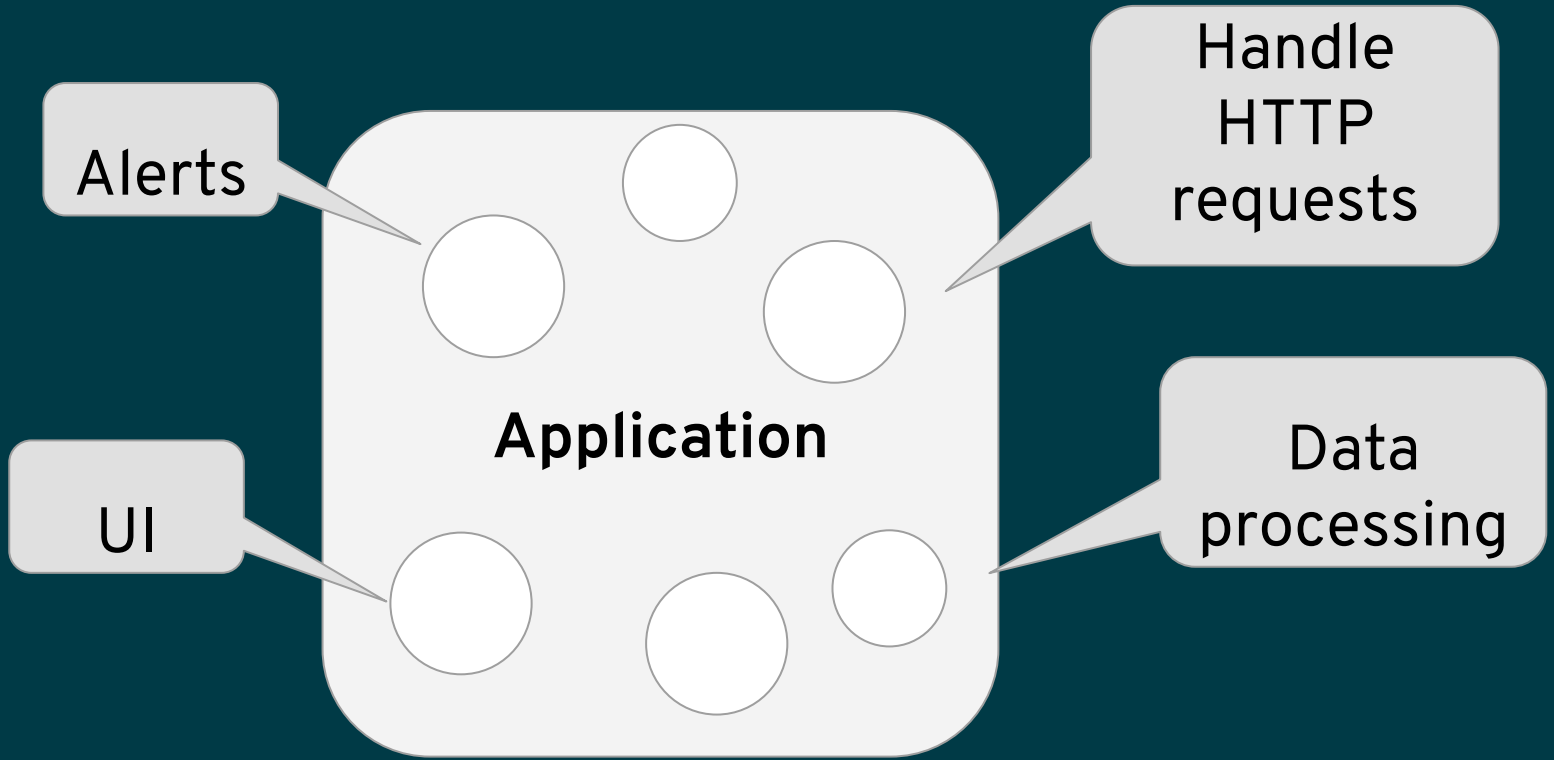

Multiple processes

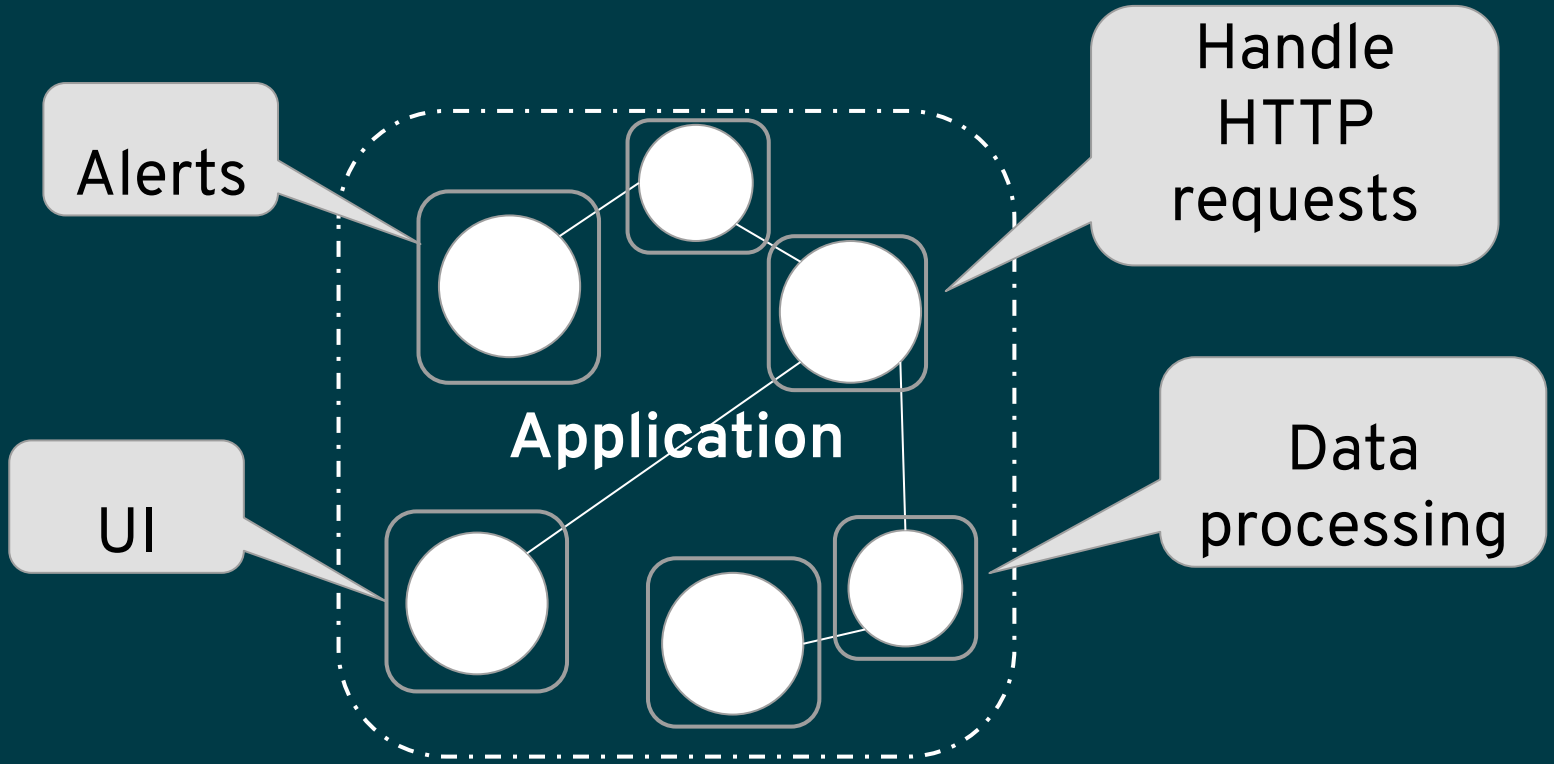

Microservices

Language agnostic

Scaled separately

Upgraded separately

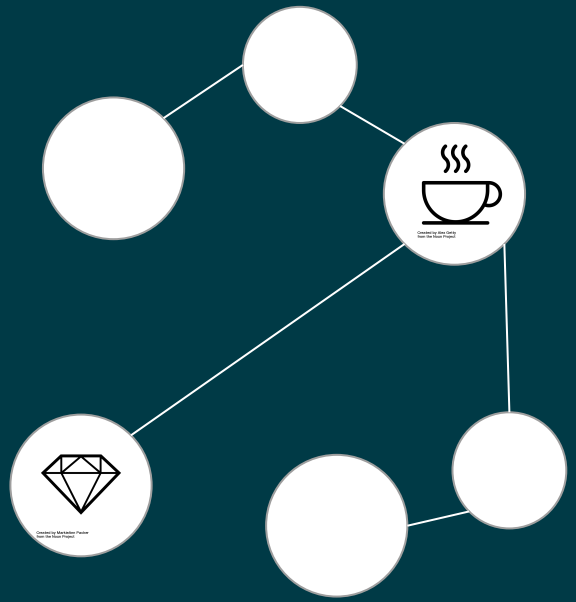

A shift in Application Packaging and Runtime

Containerizing an app

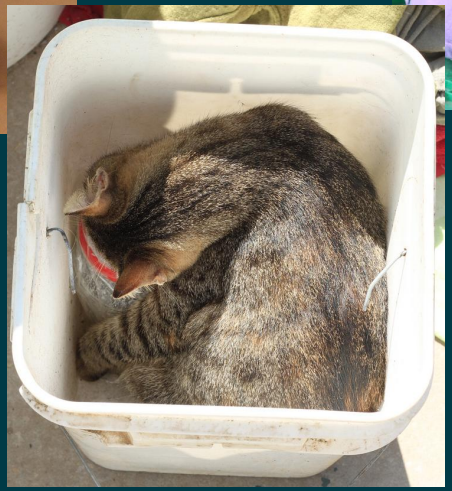

Run multiple containers

Orchestrate containers

- Run many containers on multiple hosts
- Scale - manage several instances (replicas) of the same container
- Manage a container based environment

Container orchestration platforms

Kubernetes

okd

OKD
(Openshift)

Kubernetes building blocks (some...)

- Pod - a group of one or more containers, with shared storage/network
- Deployment - manages pod definition and defines replicas of pods
- Service - an abstraction, an access point to a set of Pods
    - Sometimes called a **microservice**

Microservices - the Kubernetes way

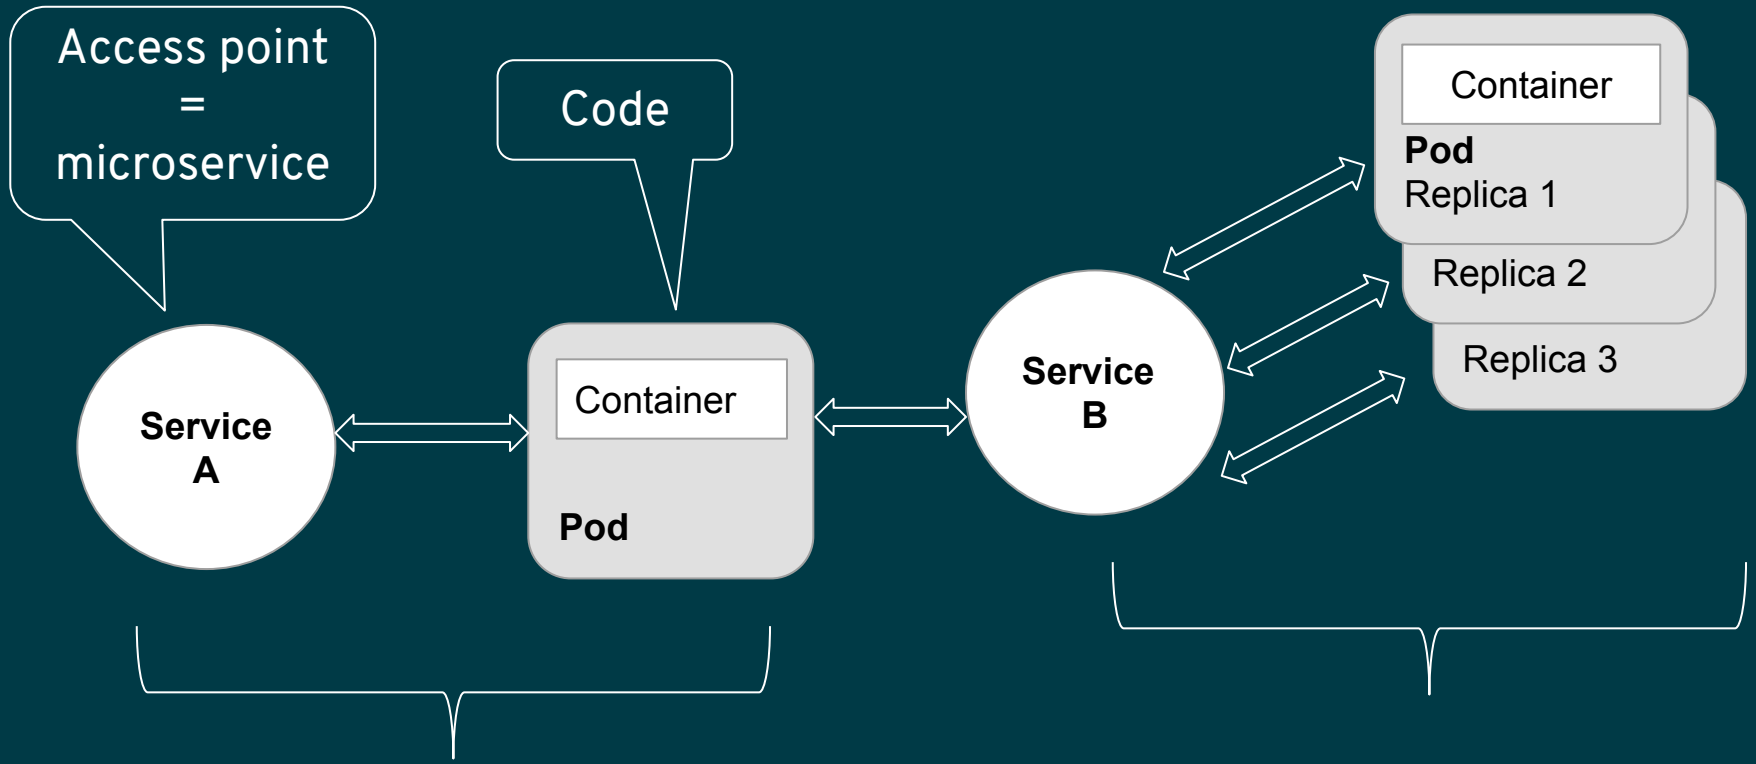

High Complexity

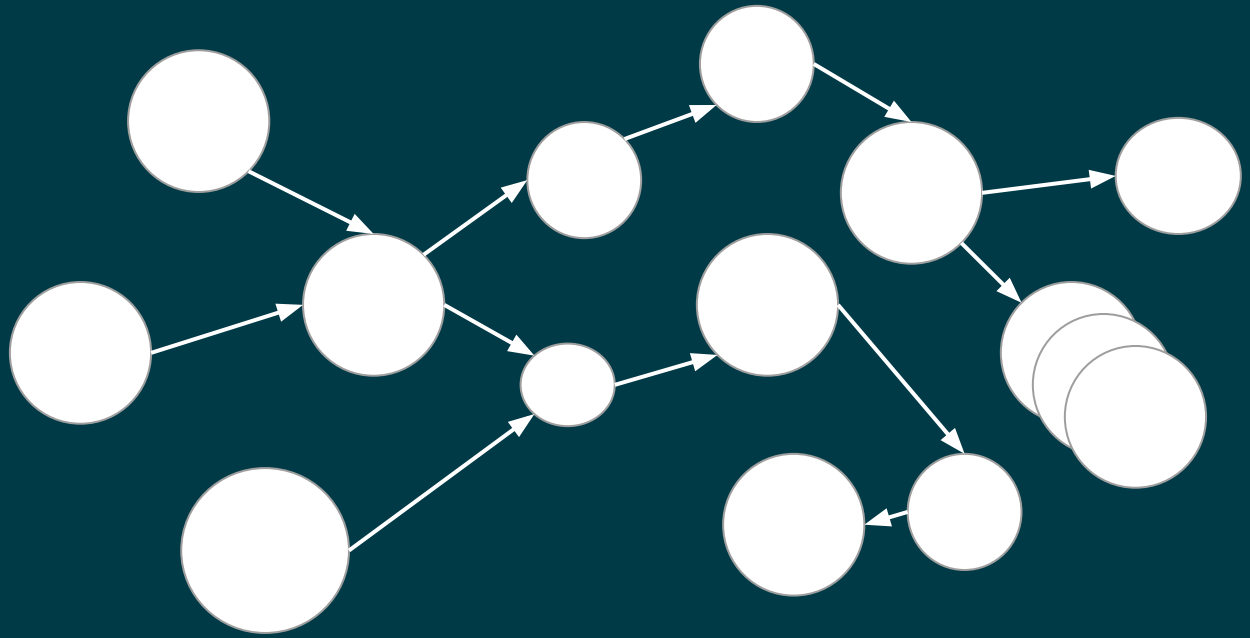

Multiple points of failure

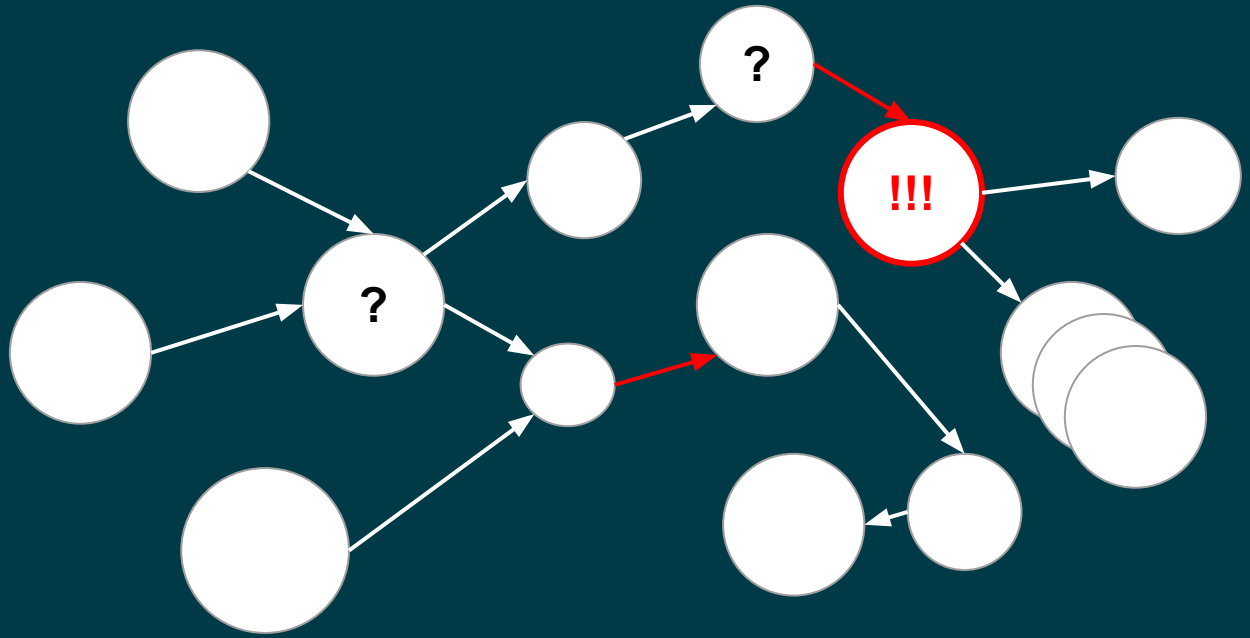

Challenges

- How are the requests routed between services?
- How do I detect failures and downtime?
- How to upgrade and test new versions of a service?
- Securing the communication

Service mesh to the rescue

What is a service mesh

- Infrastructure/framework that handles communication between services
- Often implemented as network proxies deployed alongside the microservices

Istio - Ιστίο

...

Open source service mesh

The dry facts

- Started in May 2017
- Means “sail” in Greek
- Developed in Go

Istio features

- Load balancing (HTTP, gRPC, TCP...)
- Traffic control (routing rules, retries, timeouts, fault injection, mirroring)
- Secure service-to-service communication
- Access controls (authorization)
- Metrics and traces for traffic

Important Terminology

- Workload - anything owning/controlling pods (like a Deployment) or the pods themselves
- Service - a **microservice**
- Application - *label* “app” on a pod/service
- Version - *label* “version” on a pod/service

Before Istio

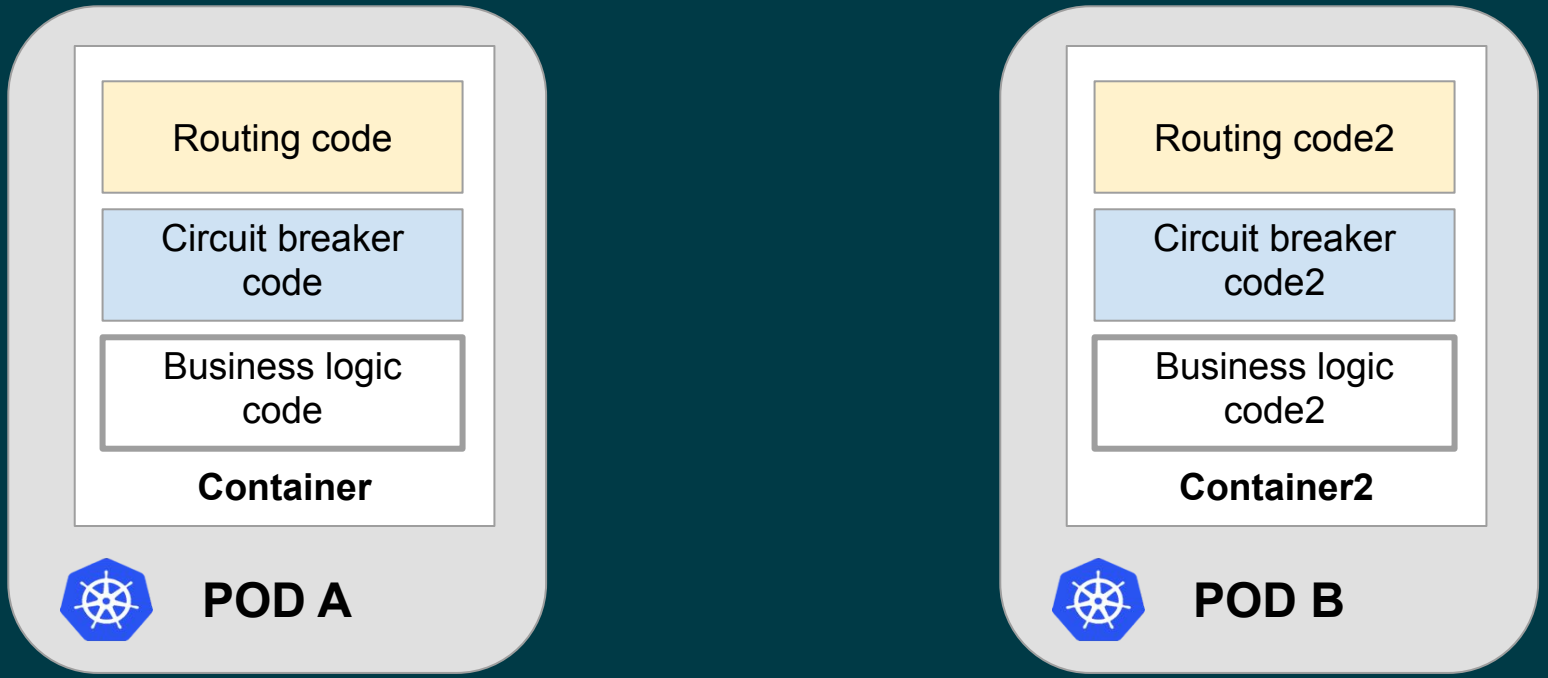

Istio

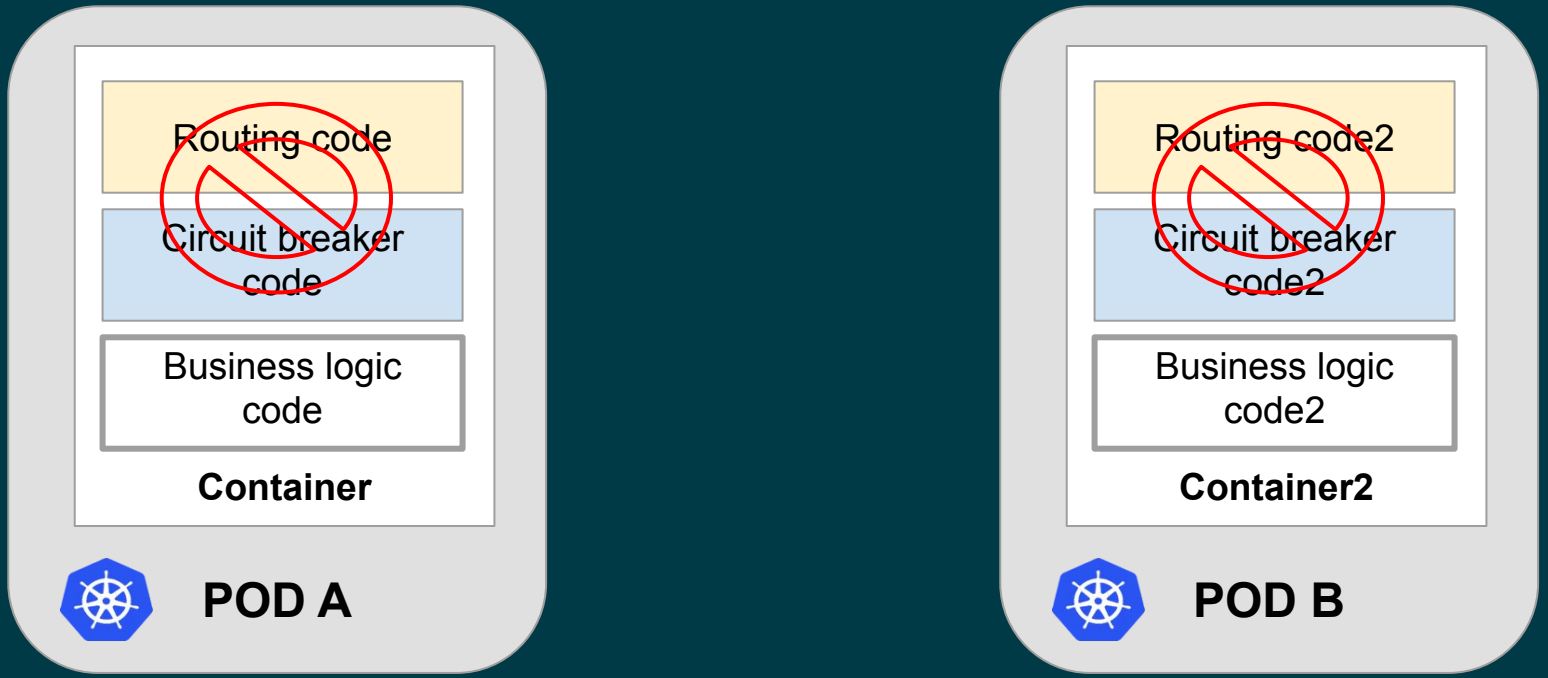
Routing code
Circuit breaker code
Business logic code
Container
POD A
Routing code2
Circuit breaker code2
Business logic code2
Container2
POD B

Sidecar Proxy

- A proxy is deployed in a container next to each instance of microservice (inside a pod)
- Container name: istio-proxy
- It is **transparent** to application code
- Envoy open source proxy is currently used 

How is the sidecar injected?

- Manually
- Automatically injected to pod on creation
  - *kubectl label namespace default istio-injection=enabled*
  - Mutating Admission Webhook is used for sidecar injection
  - Actually… 2 containers are injected: istio-init and istio-proxy

Istio architecture

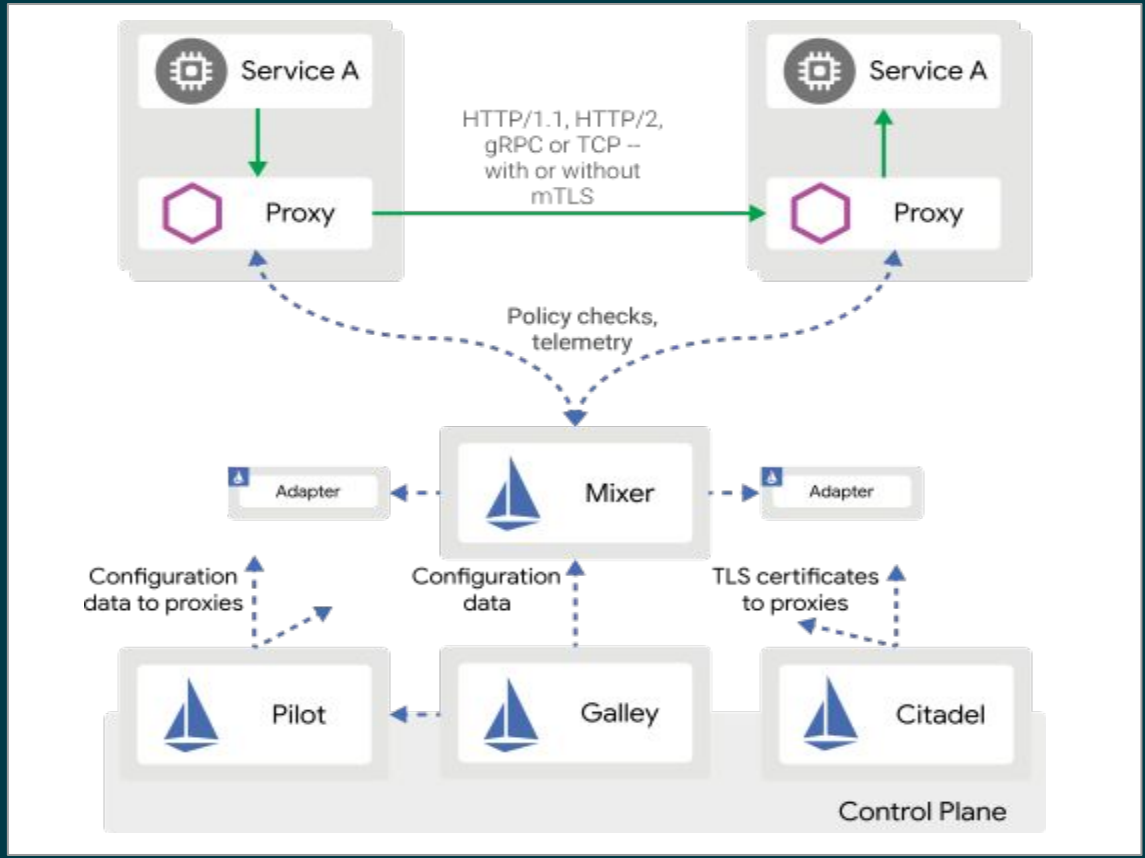

Sidecar Proxy in Istio and Kubernetes

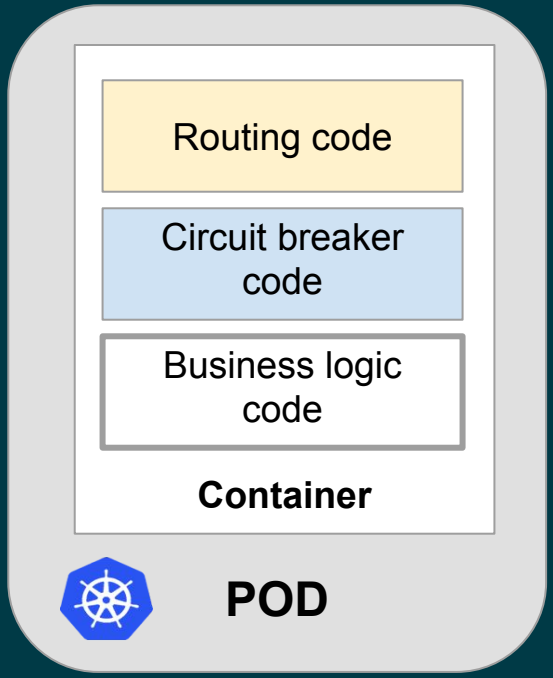

Before Istio, no sidecar

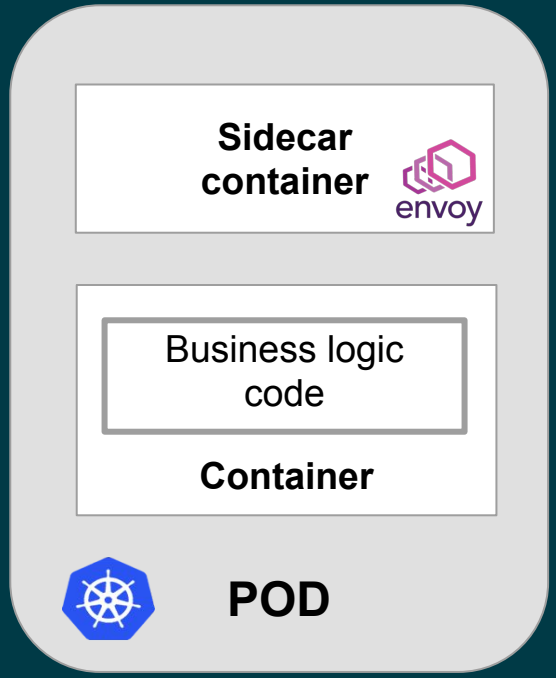

With sidecar

With Istio - sidecar intercepts all traffic

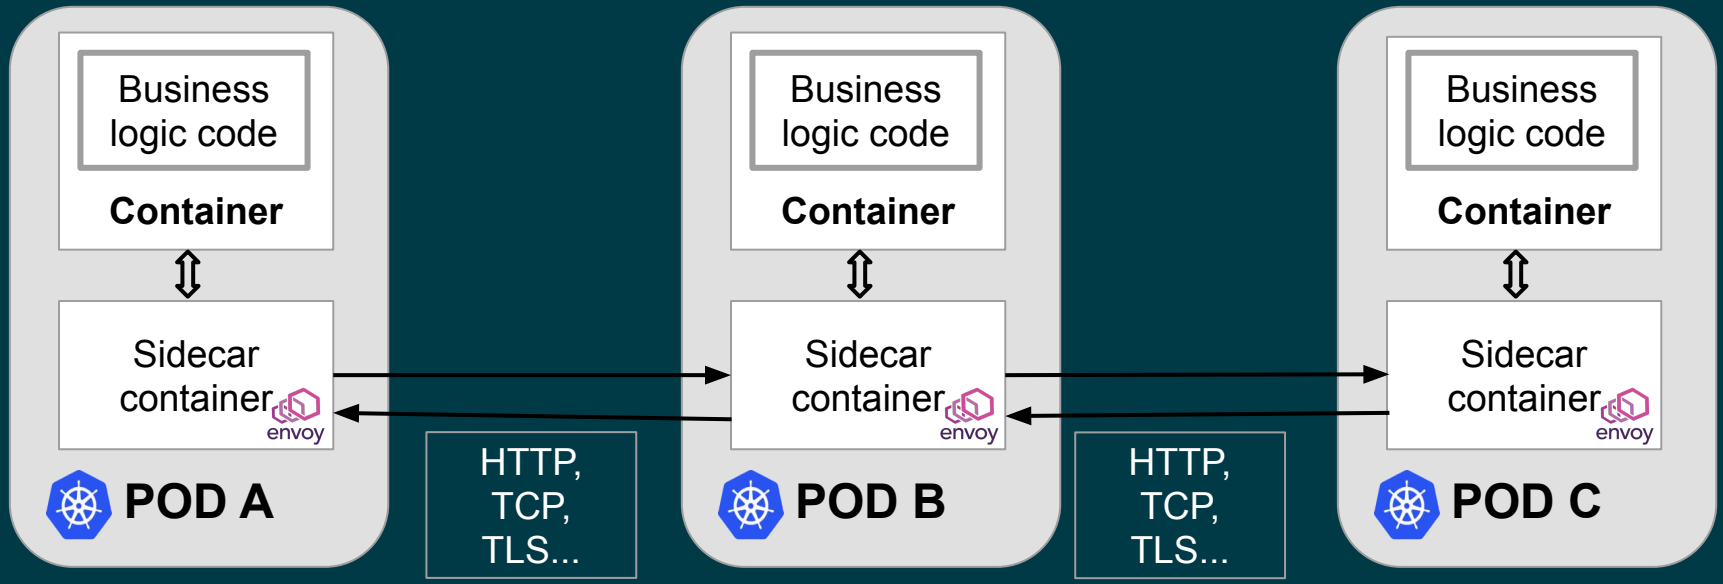

Configuration is transparent to the services and not part of the code

Istio routing in Kubernetes

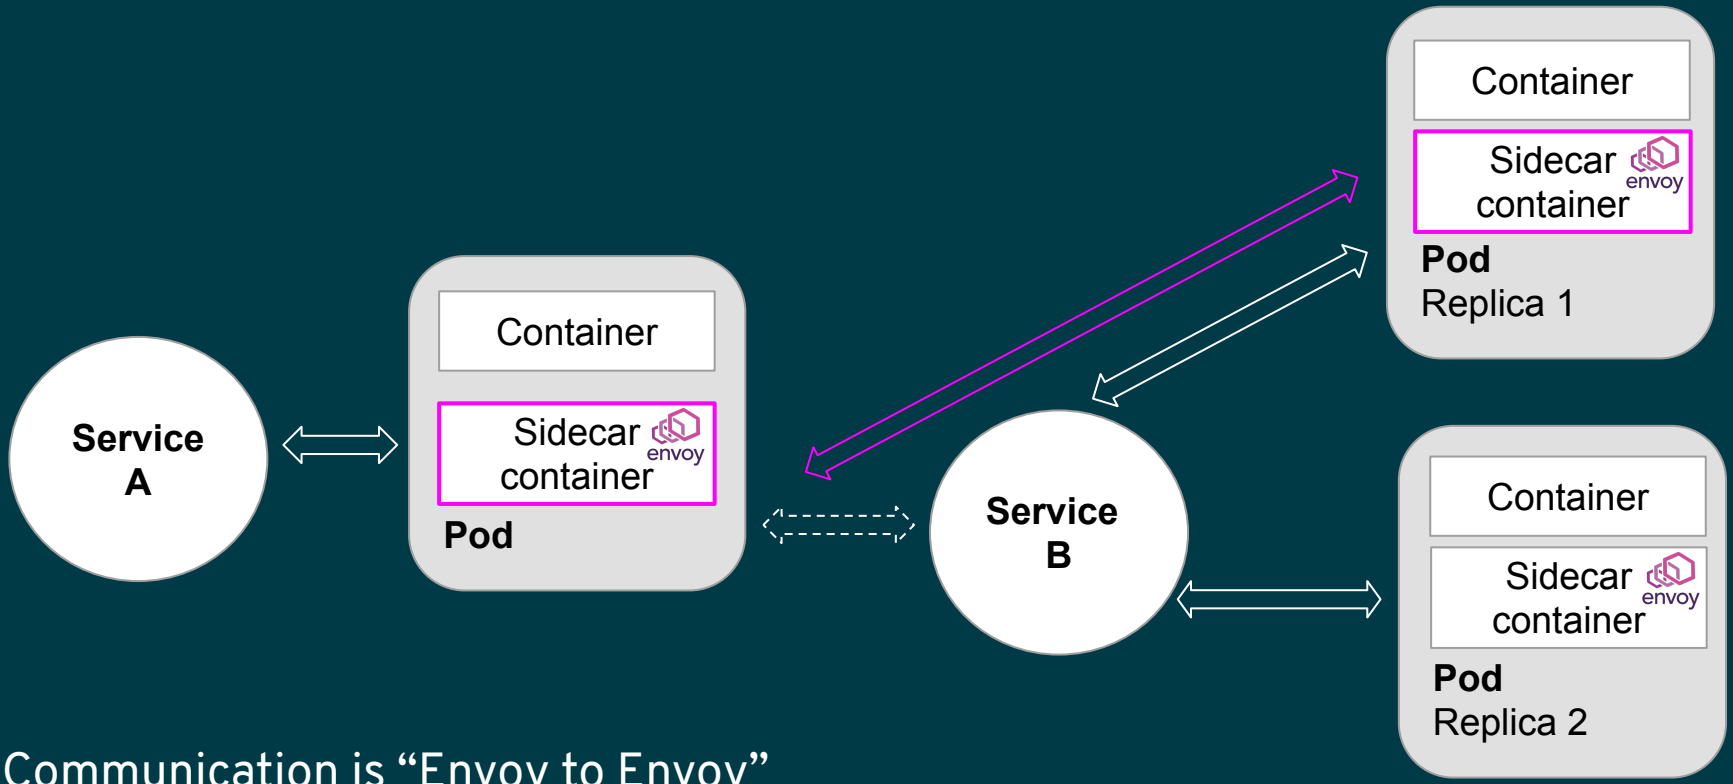

Communication is “Envoy to Envoy”
bypassing the Kubernetes Service

Different routing scenarios

- A/B testing
- Traffic shifting
  - Canary deployment (an example of traffic shifting)
- Mirroring traffic

Weighted Routing with Istio - A/B

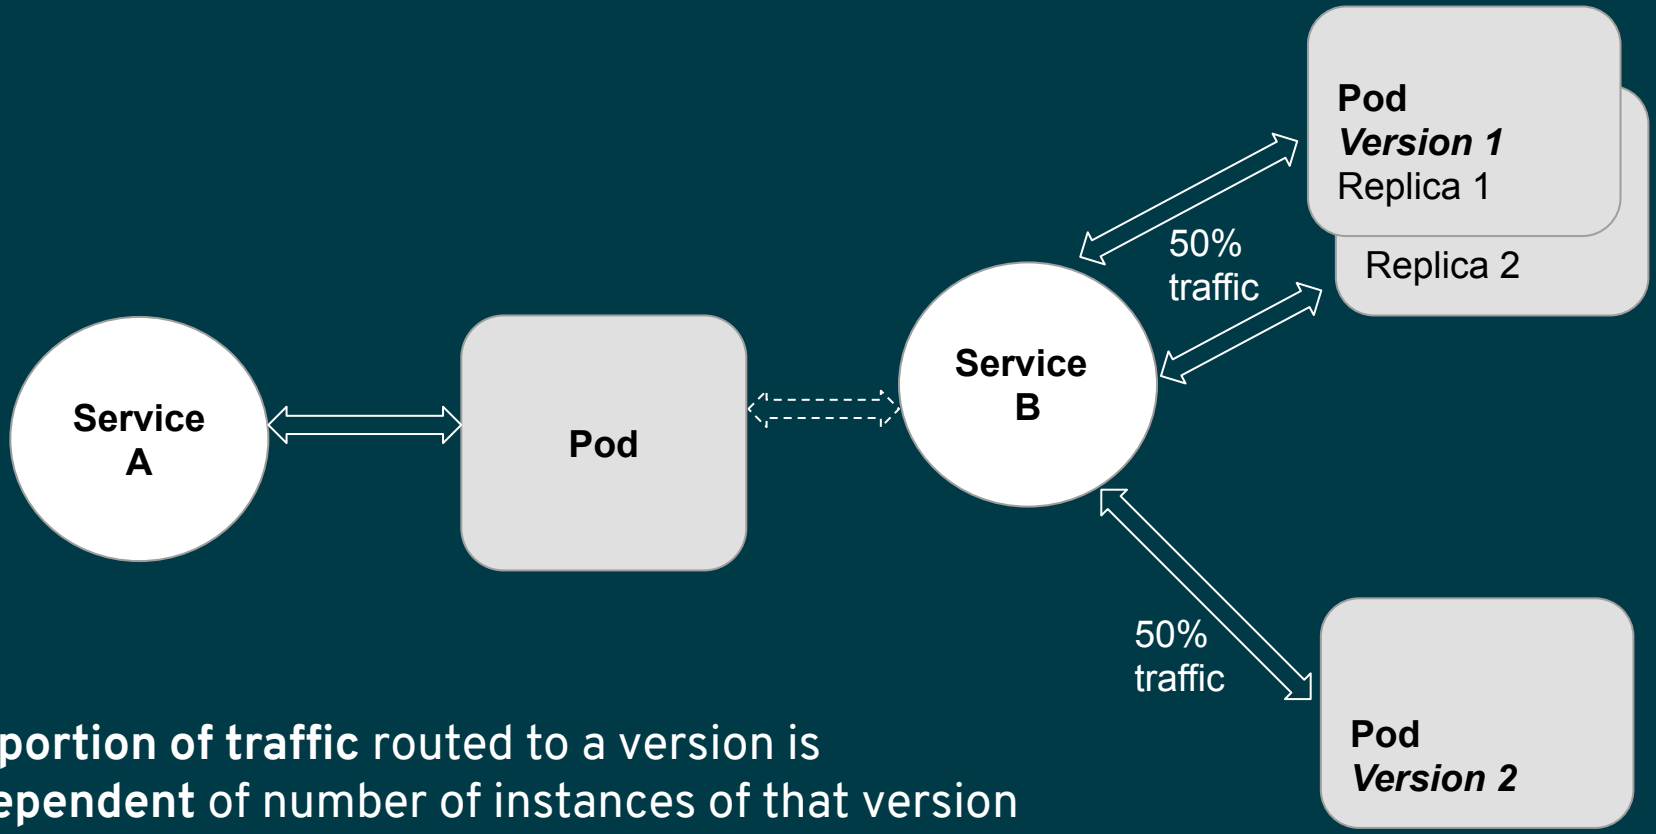

Proportion of traffic routed to a version is **independent** of number of instances of that version

Weighted Routing - Canary

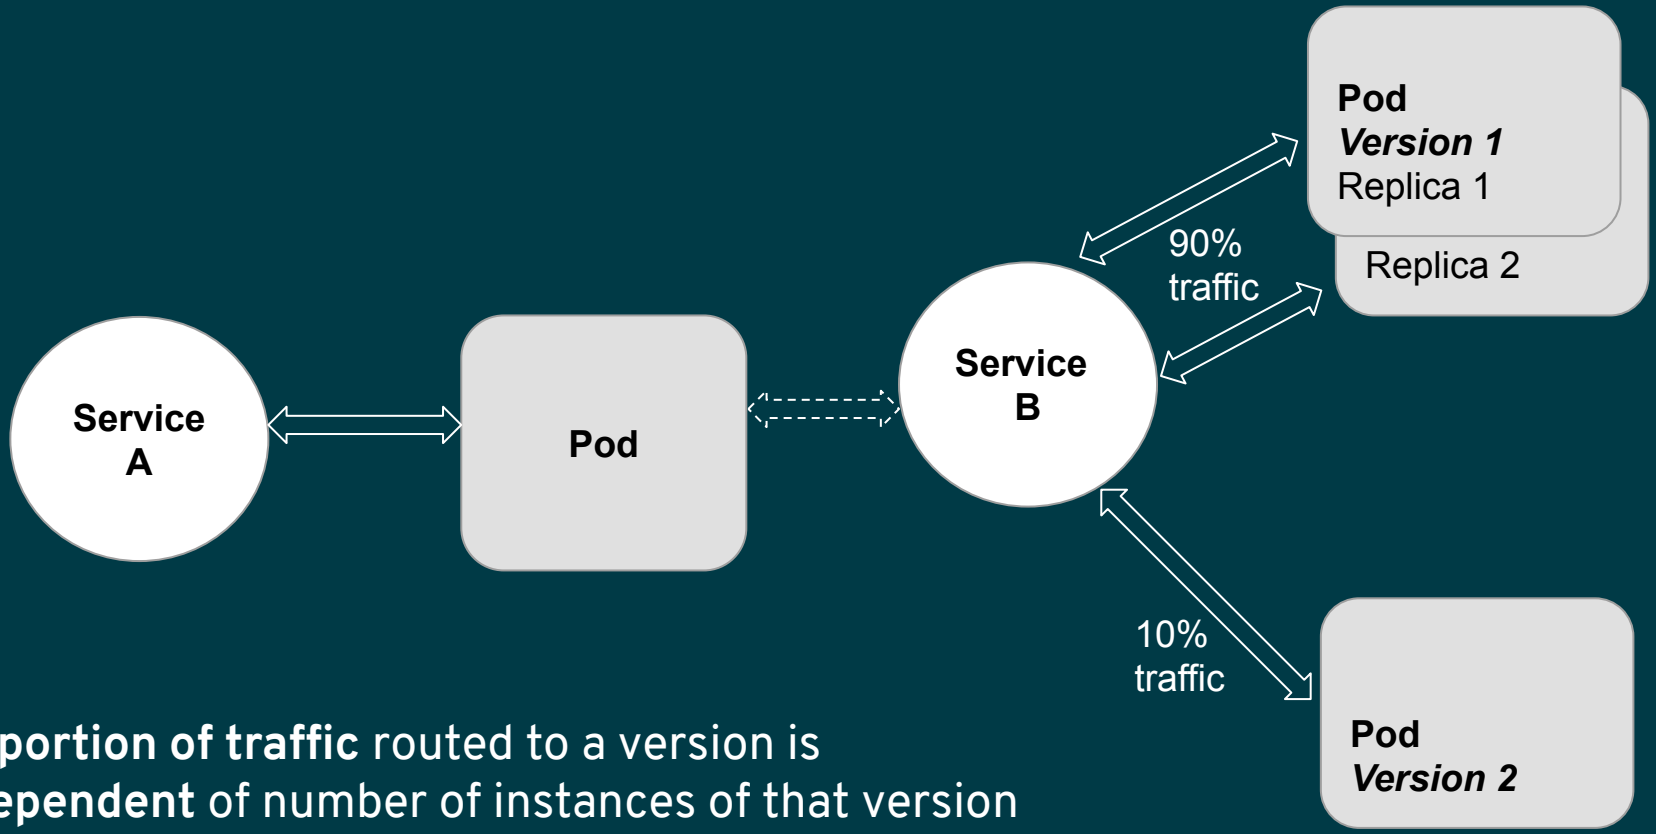

Proportion of traffic routed to a version is **independent** of number of instances of that version

Matching Routing with Istio

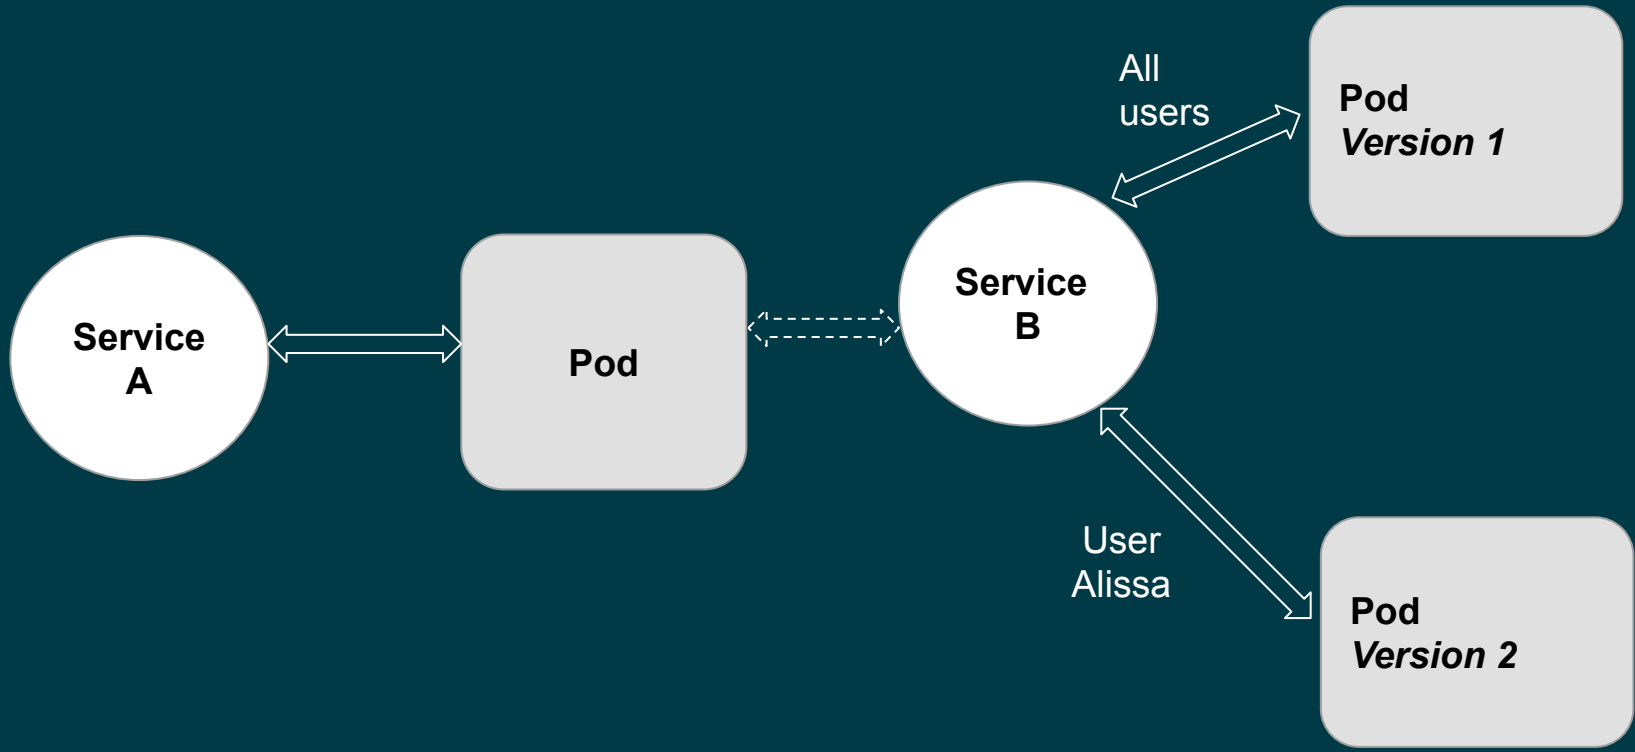

Mirroring traffic

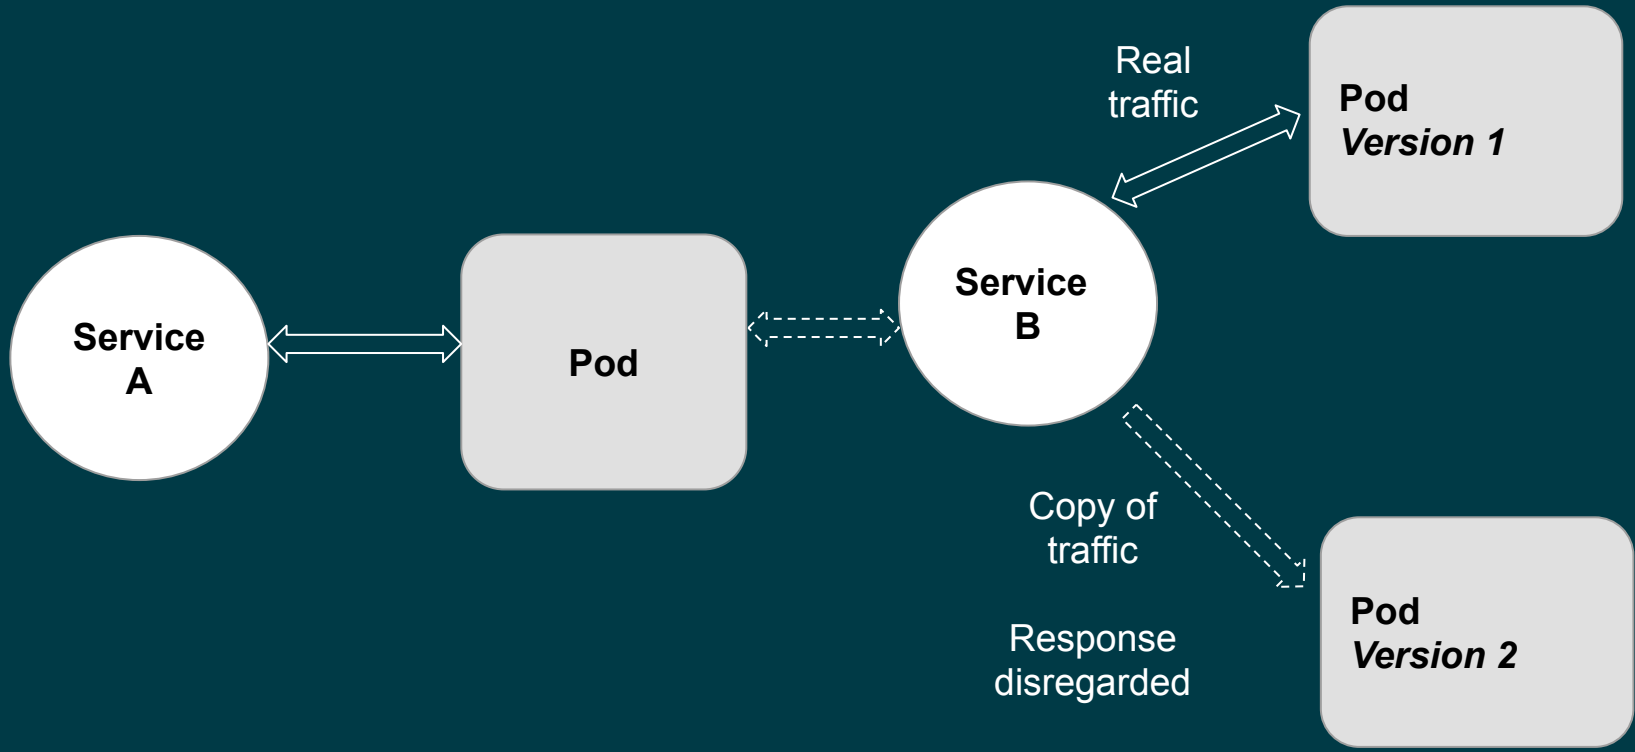

"Anything that can go wrong will go wrong"

(Murphy’s law)

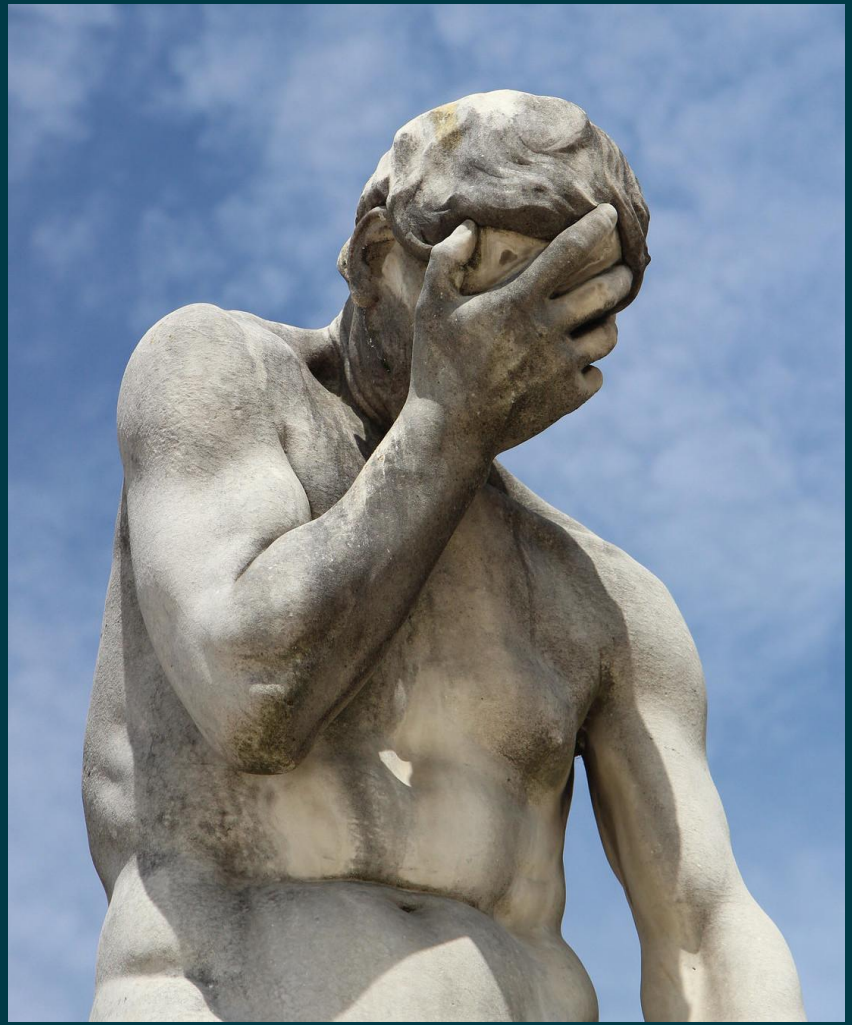


KEEP
CALM
AND
DO
CHAOS
ENGINEERING

Chaos engineering with Istio

- Inject delays
  - Simulate network latency
  - Simulate an overloaded service

- Define aborts (Inject Errors)
  - Simulate failure in a service (return a predefined HTTP Error)
  - A good alternative for a manual shutdown or “scale to zero”

Inject delay

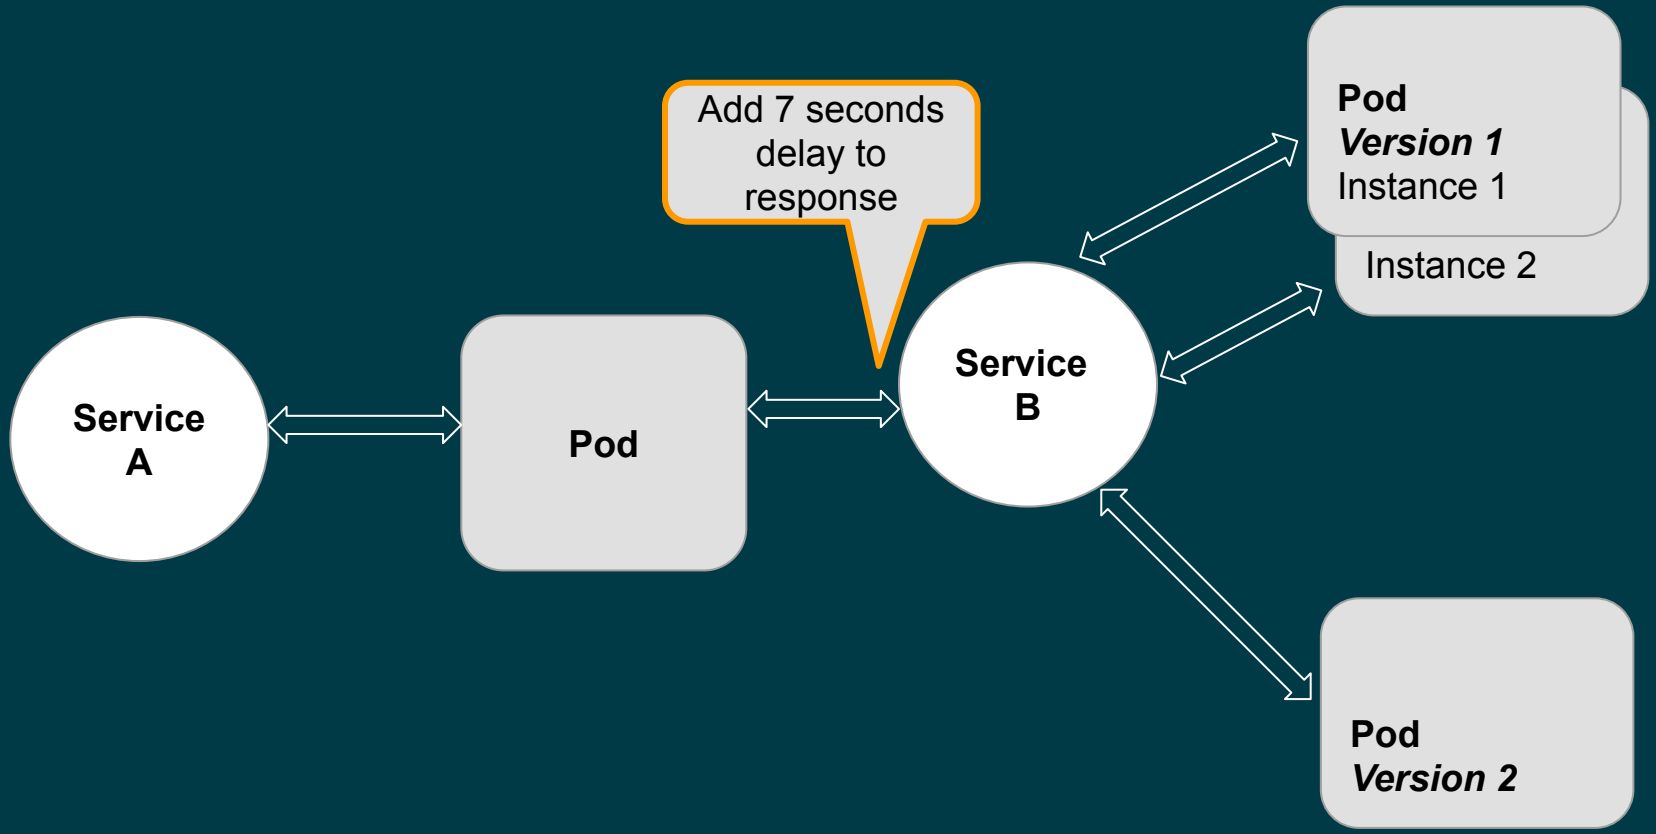

Inject Error

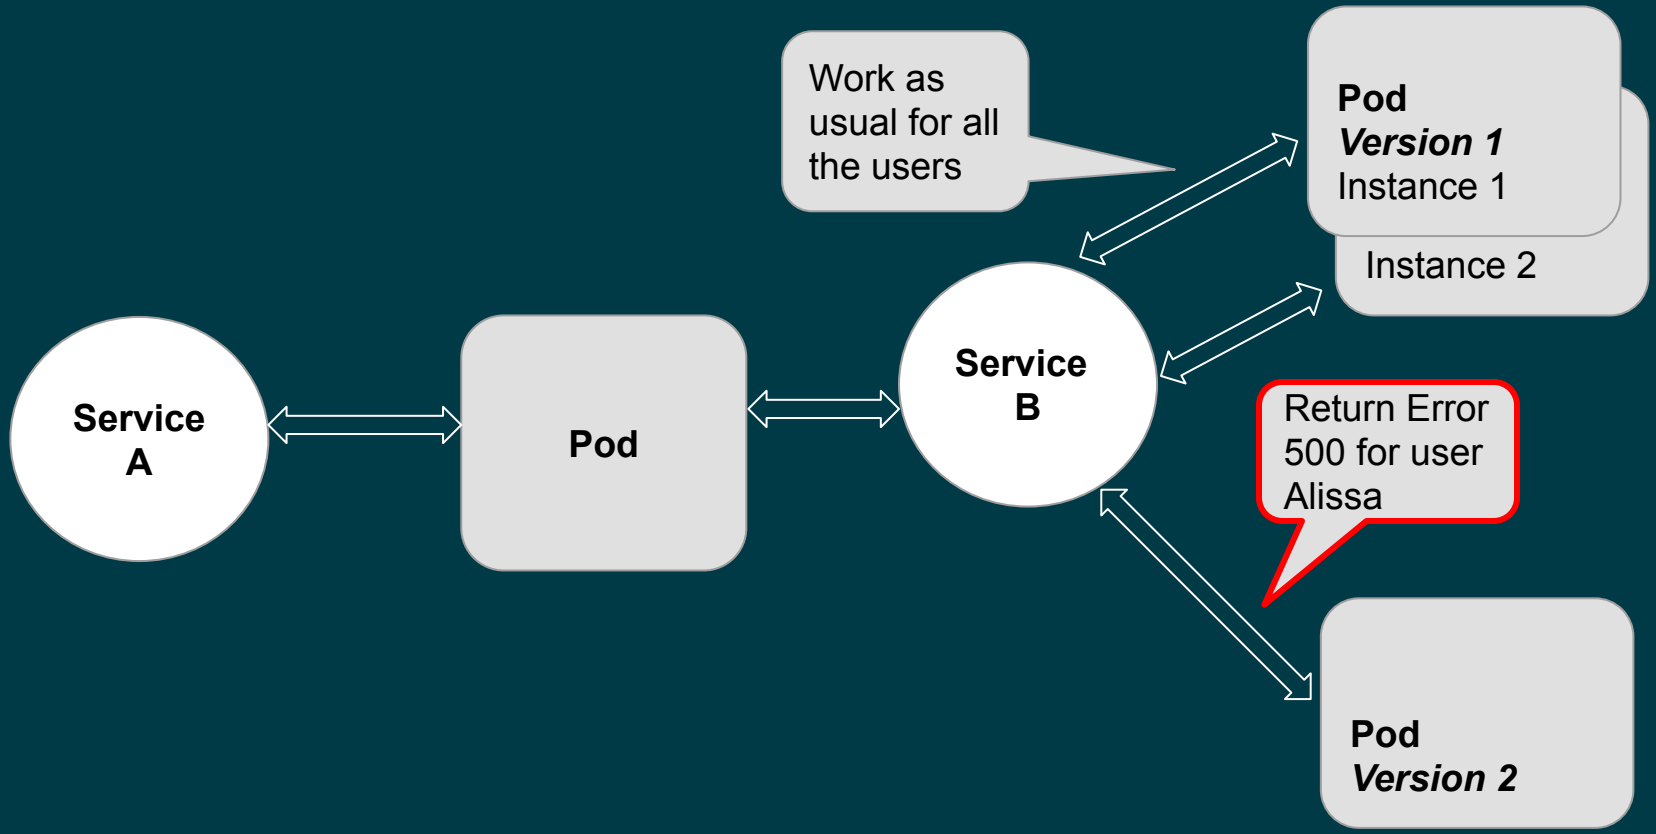

Circuit breaker

- Set a connection pool to limit connections and requests
- **Example:** “Set a connection pool of 100 connections with no more than 10 req/connection to service A”

Outlier detection

- Classify instances as healthy/unhealthy
- Eject unhealthy instances for a defined timeframe which can be increased over time
- **Example**: “Scan all pods every 5 mins, any instance that fails 7 consecutive times with 5XX error code will be ejected for 15 minutes.”

Authorization and Authentication

- Authentication
  - End user authentication (JSON Web Token (JWT))
  - Service to service authentication (mutual TLS)
    - Permissive mode is possible for flexible migration
- Authorization
  - Can service <A> send <this request> to service <B> ?
  - Roles are visible across namespaces
  - ServiceRole and ServiceRoleBinding

Security

- Defining a Gateway ingress/egress to enable traffic in/out of mesh
- Citadel monitors service accounts creation and creates a certificate for them
  - Certificates only in memory, sent to Envoy via SDS API
- mTLS can be defined on multiple levels
  - Client and server exchange certificates, 2 way
  - All mesh, specific service, etc.

Configuration objects

- VirtualService != Kubernetes service
  - Rules for how requests to a service are routed within service mesh
  - Routing logic, load weighting, chaos injection
- DestinationRule
  - Configures policies to be applied to a request **after** VirtualService routing has occurred
  - Load balancer, circuit breaker
- MeshPolicy, Gateway, ServiceEntry and more...

Configuration Yaml example

All Istio objects are CRD (CustomResource Definition)

```yaml
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
  - route:
    - destination:
        host: reviews
        subset: v1
      weight: 50
    - destination:
        host: reviews
        subset: v2
      weight: 25
    - destination:
        host: reviews
        subset: v3
      weight: 25
```

New set of challenges

- How many versions exist for service A?
- Is there any traffic **now**?
- Is **routing configured** for service B?
- Is my configuration **valid**?
- Is security **on**?
- Is the app **healthy**?

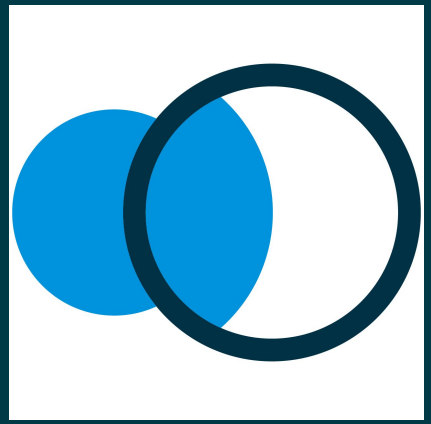

Kiali - Κιάλι

...

Open source
Istio service mesh observability

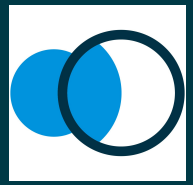

Dry facts

- Started in January 2018
- Means “spyglass” or “monocular” in Greek
- Developed in Go and React

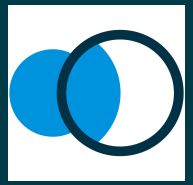

Kiali Features

- Visualize mesh connections and traffic
- Service and application health
- Configure routing via UI
- Validate Istio configurations
- View metrics, traces and logs
- Visualize security configuration

A picture is worth a thousand yamls

Demos based on Bookinfo example

Let’s see Kiali in action

- Mesh visualization
- Fault Injection
- Configuration Validation
- Configure routing rules
- Tracing
- Traffic stats

Bookinfo example

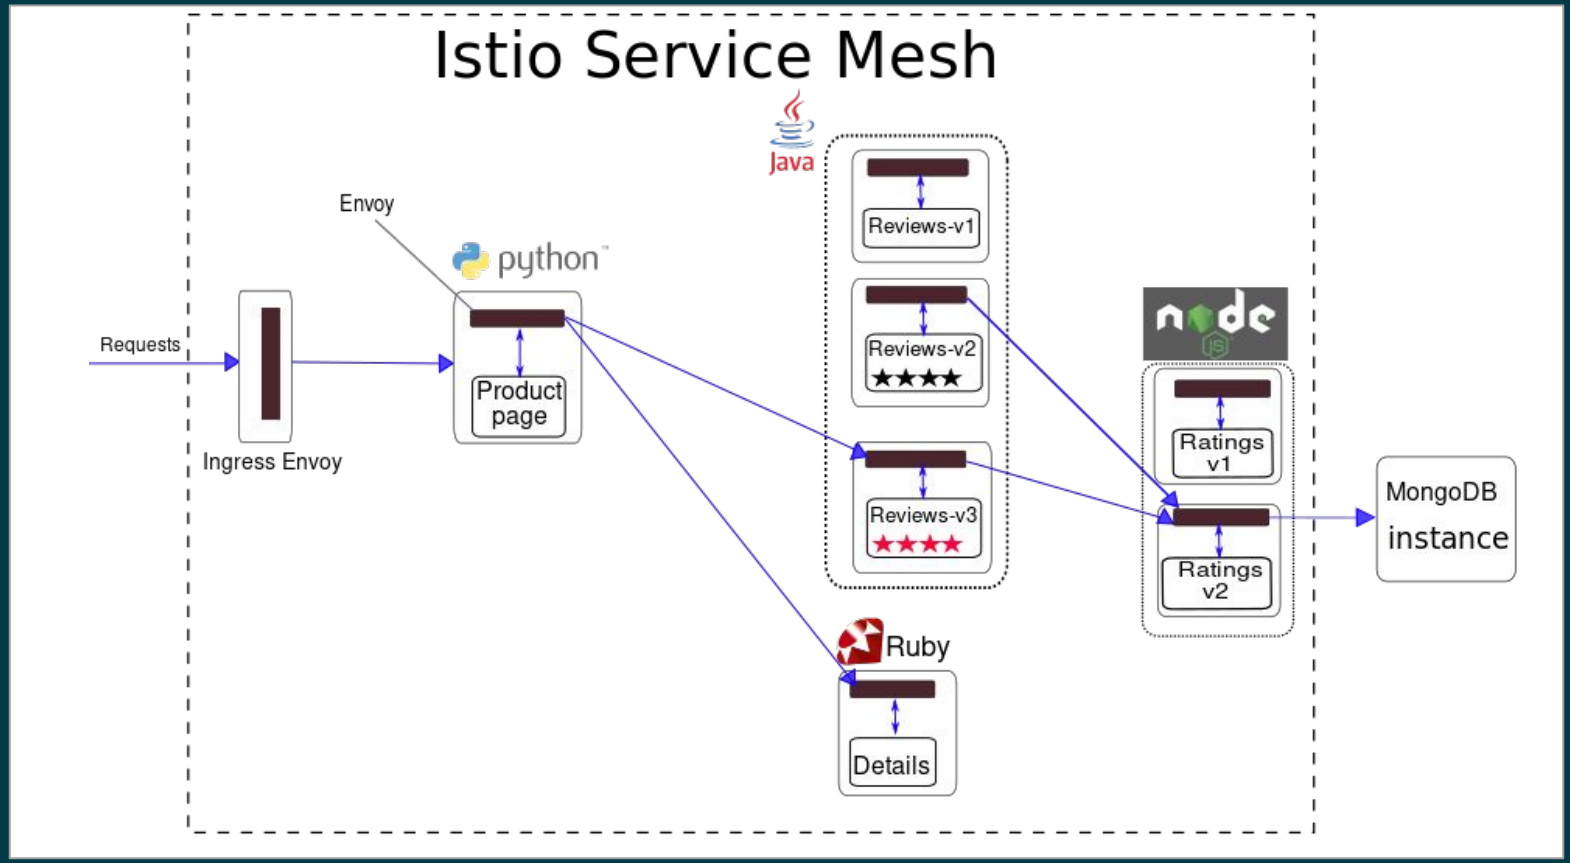

Bookinfo on Kiali

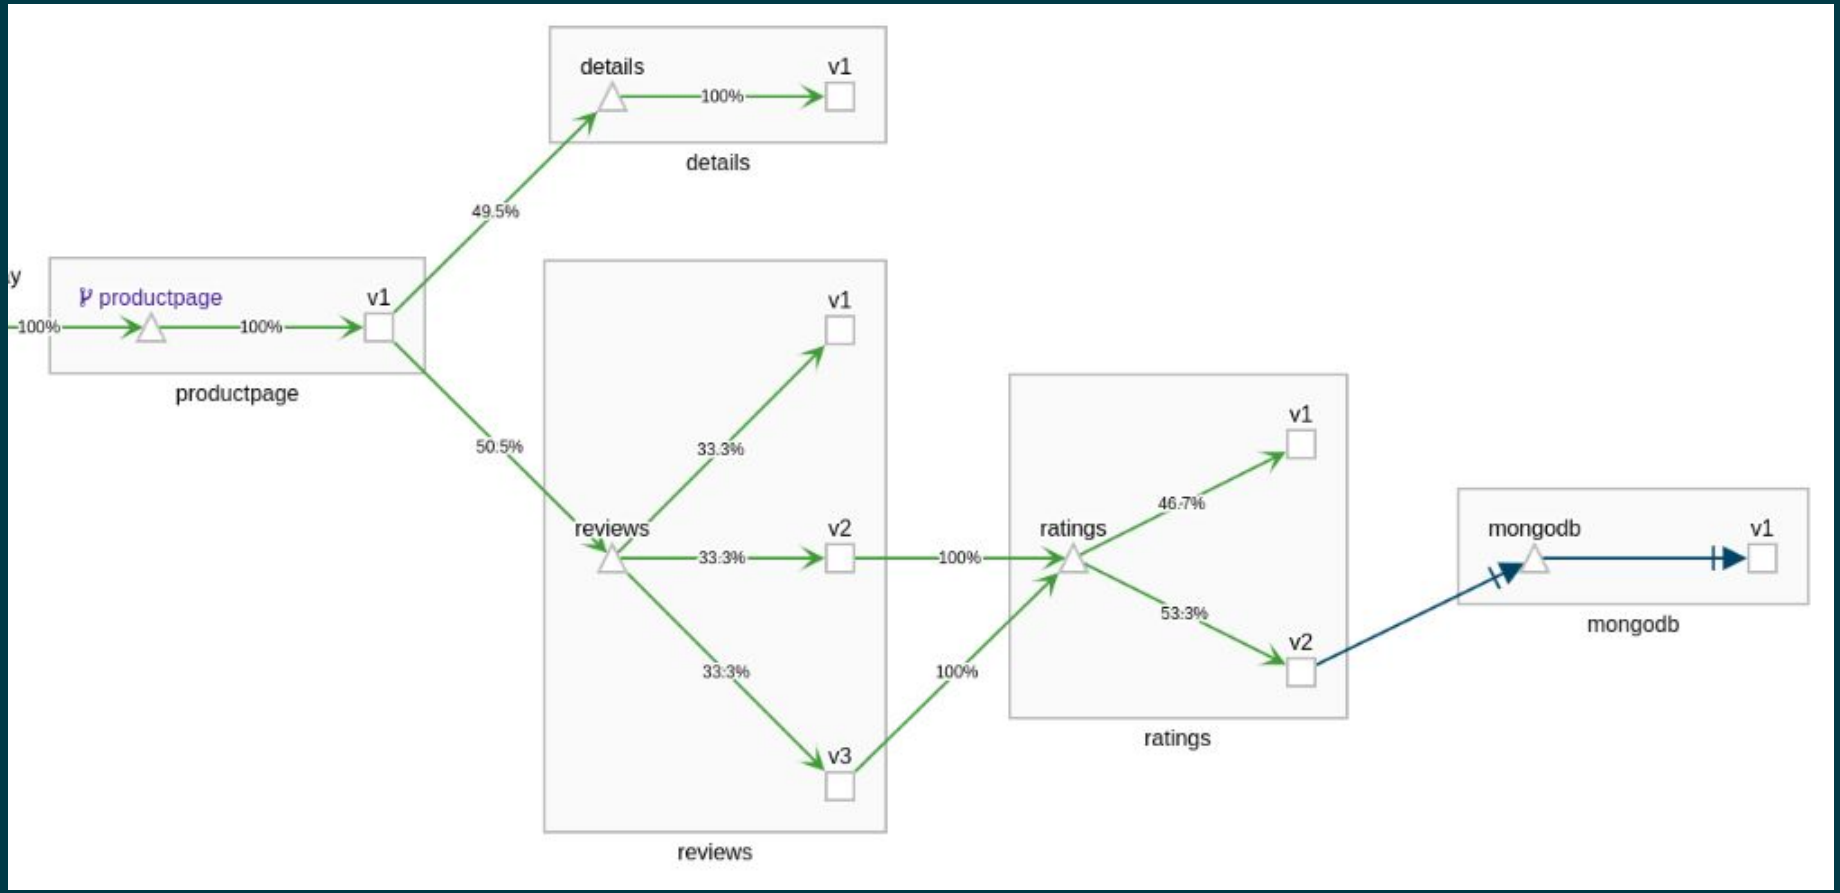

Kiali Features

Overview page

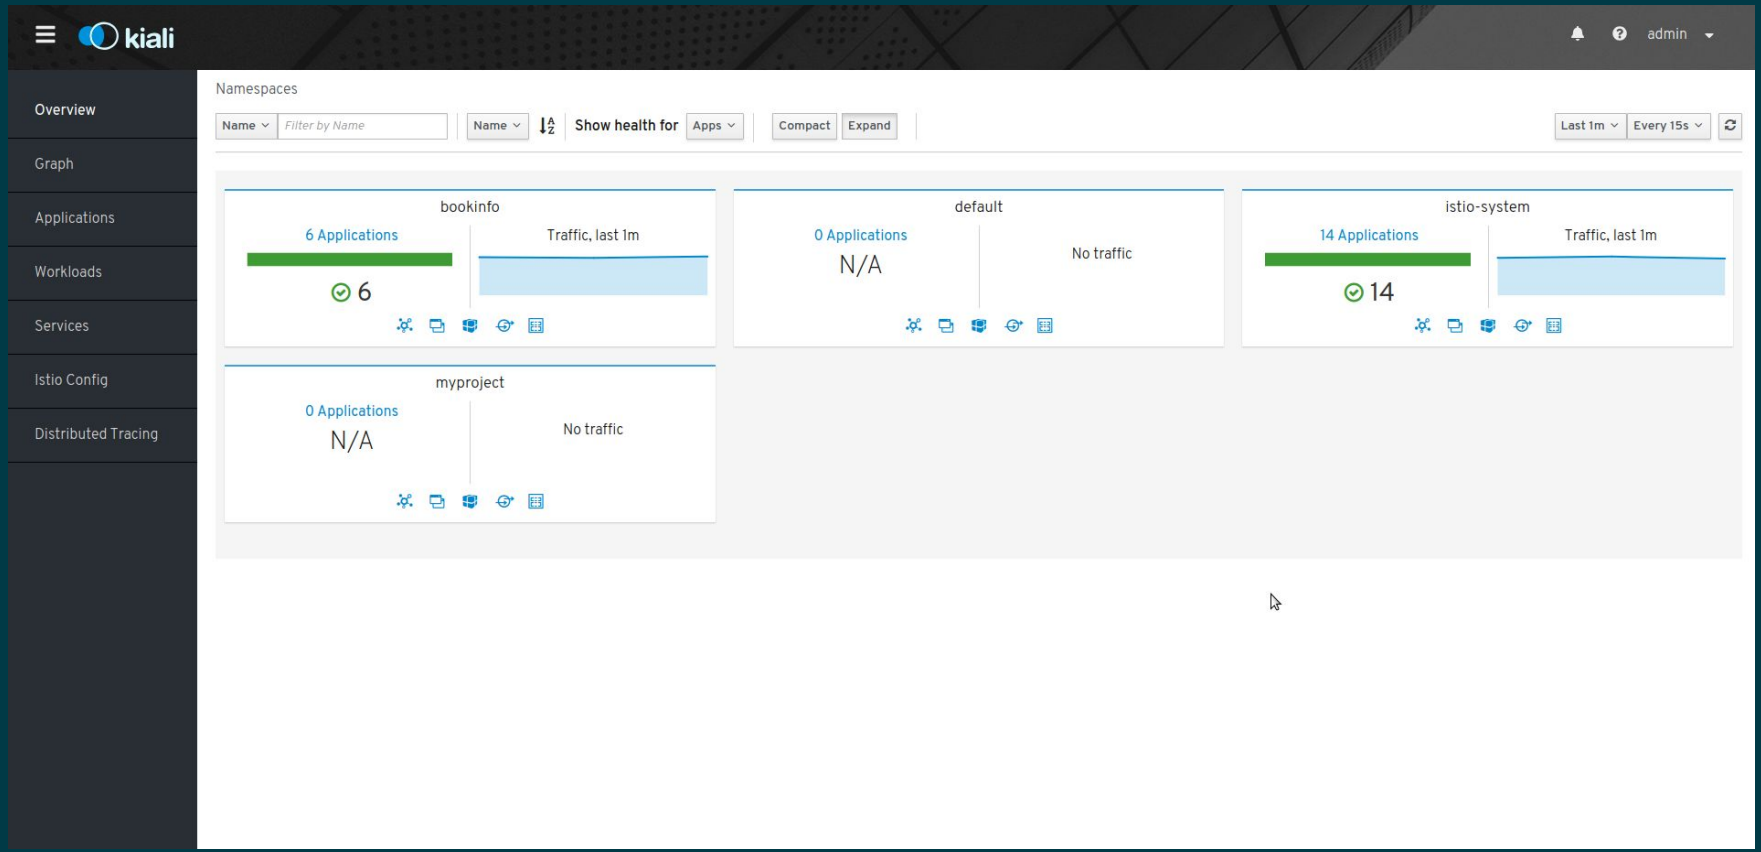

Mesh Topology Graph

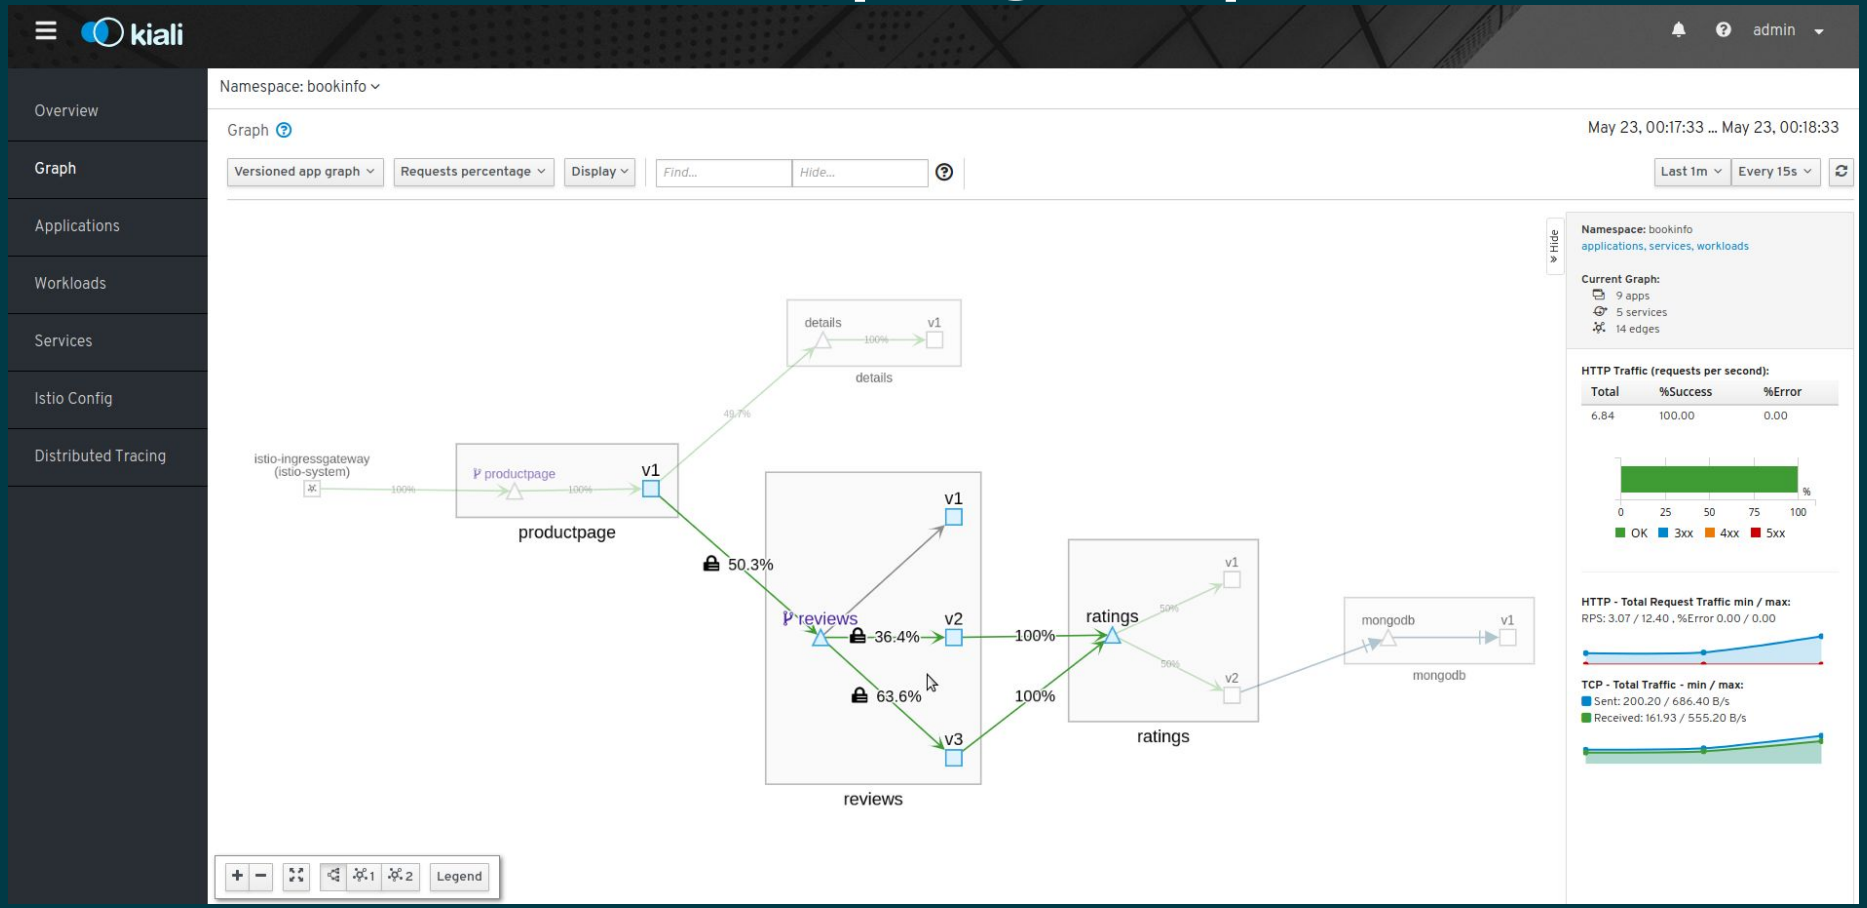

Hide and Seek

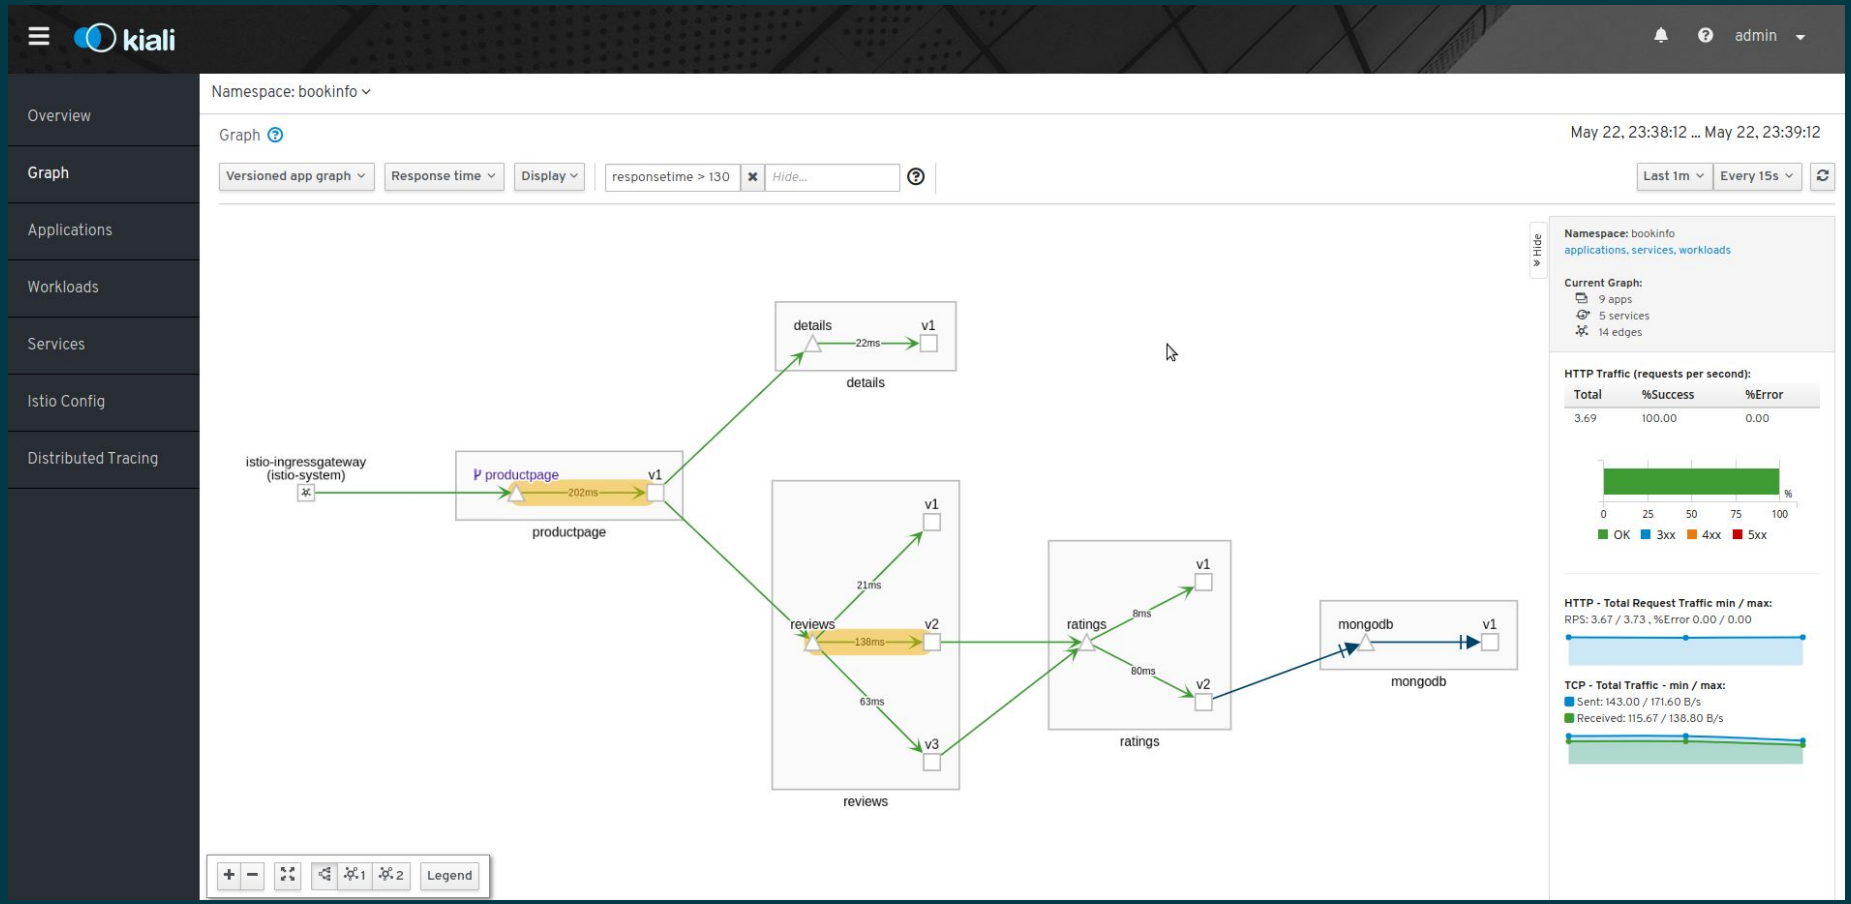

Details Page

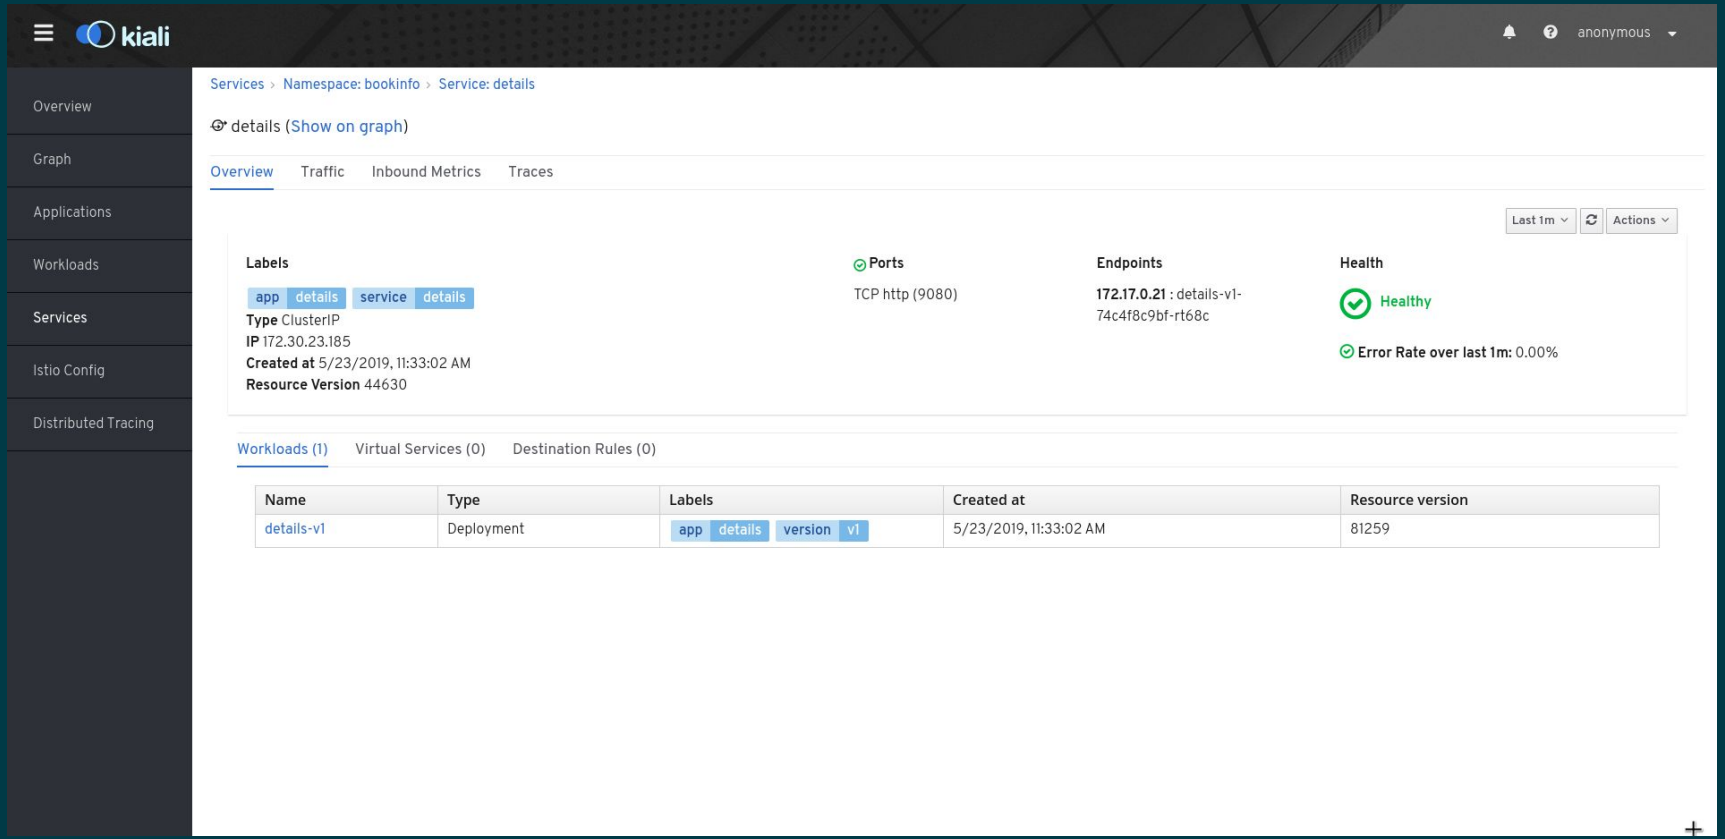

Viewing Logs

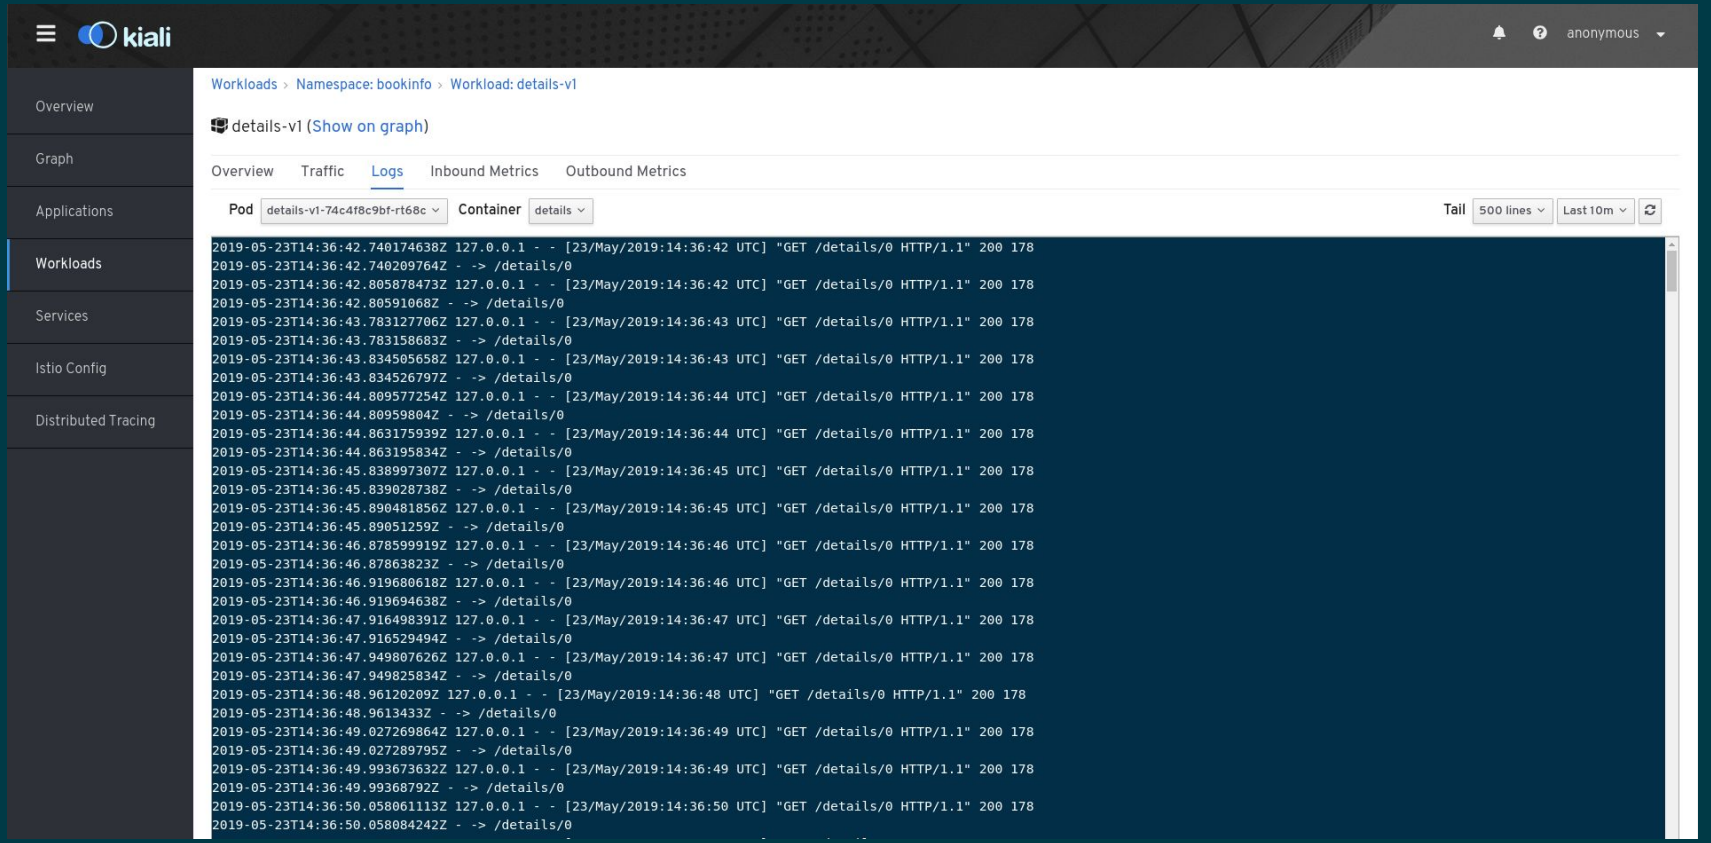

Runtime metric dashboards

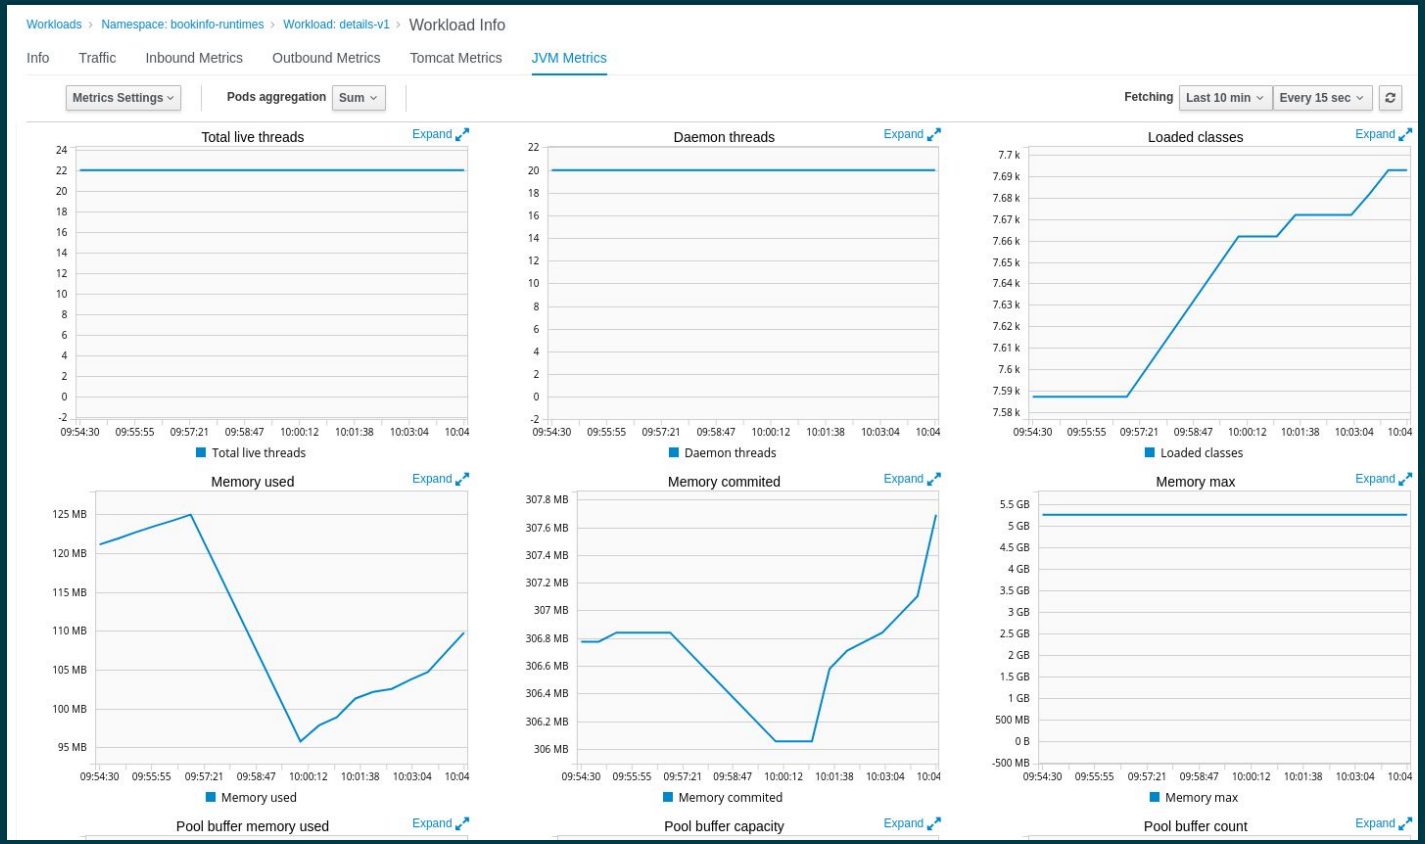

Weighted Routing

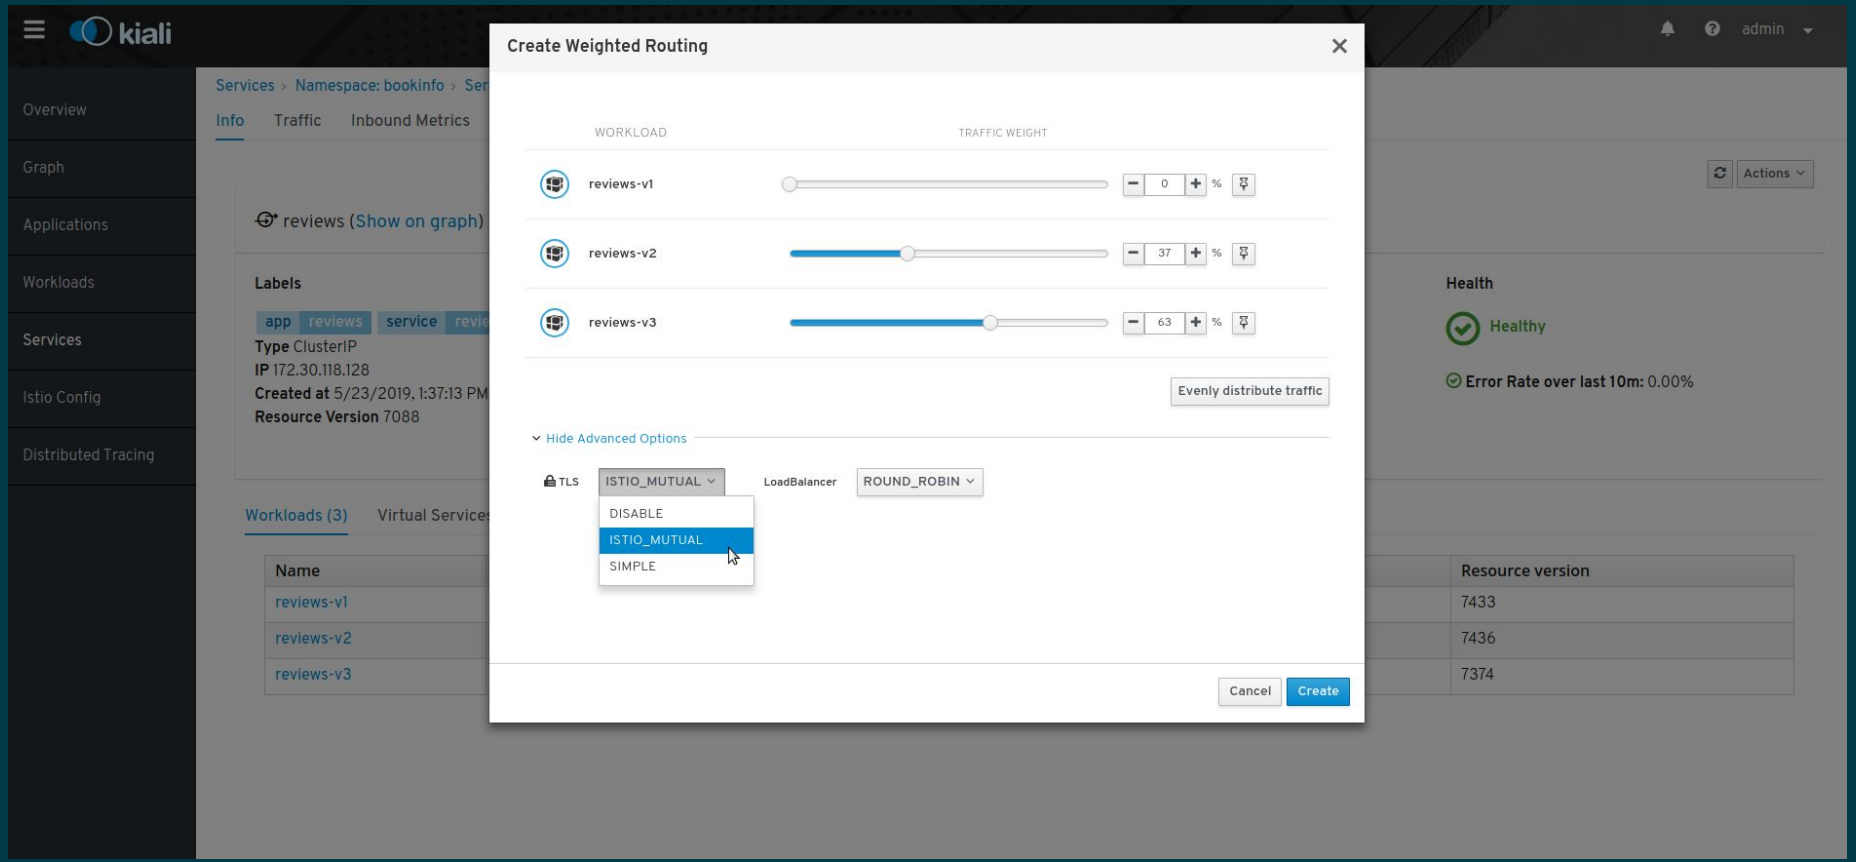

Configuration validations

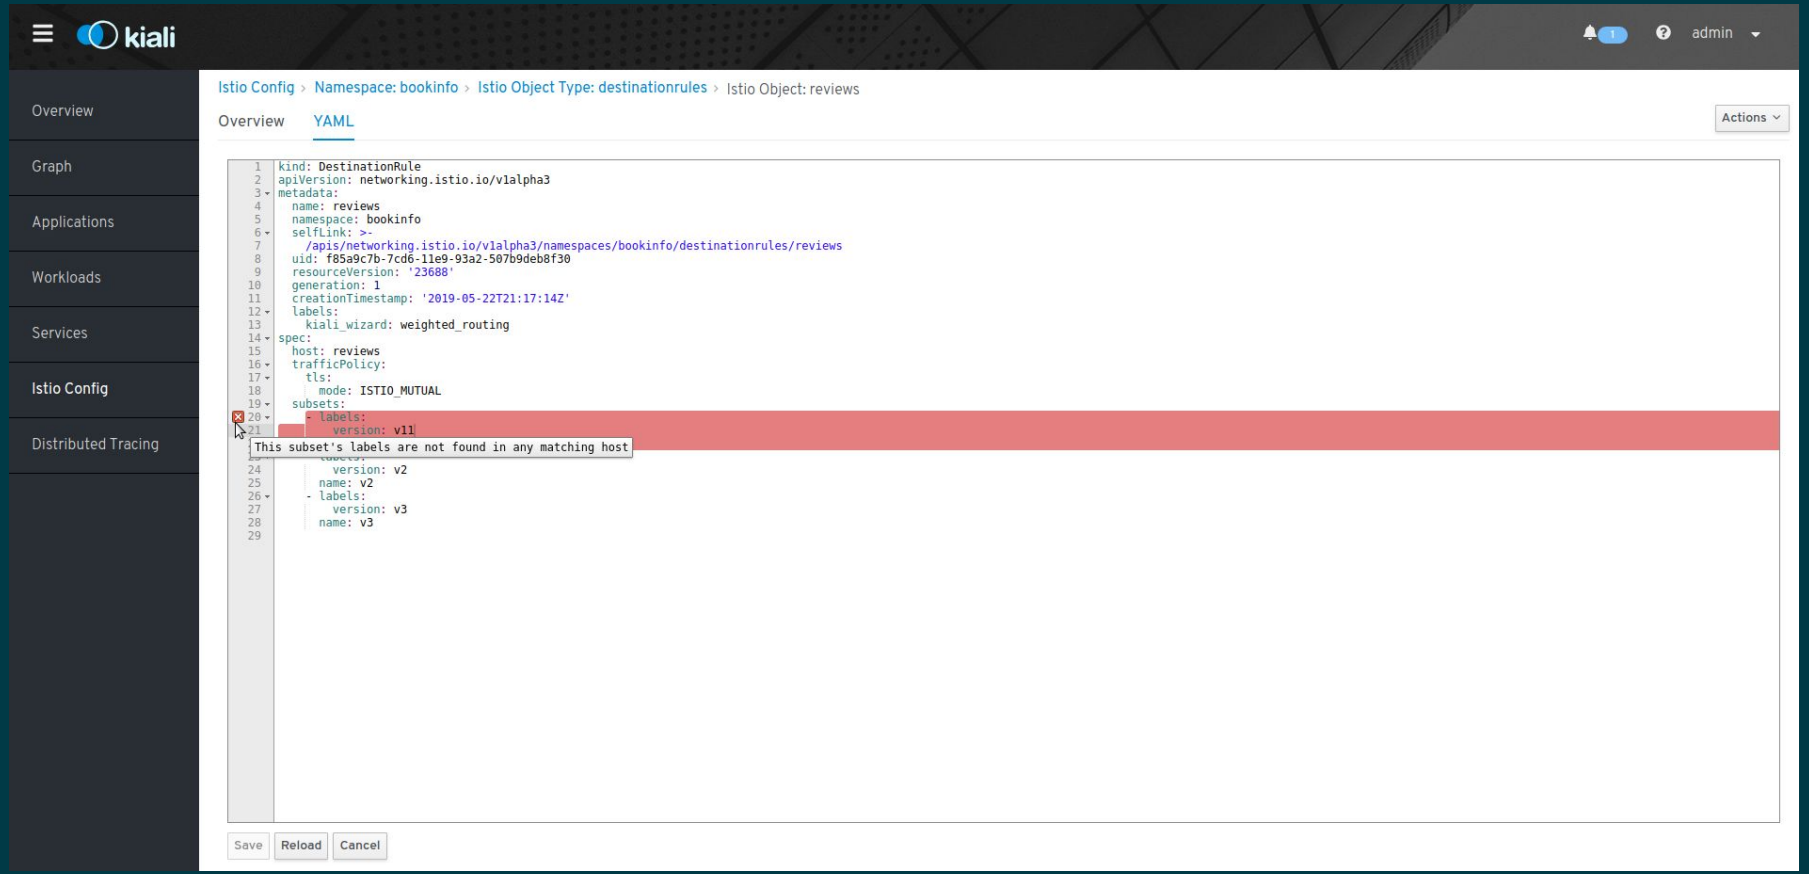

Tracing (integration with Jaeger)

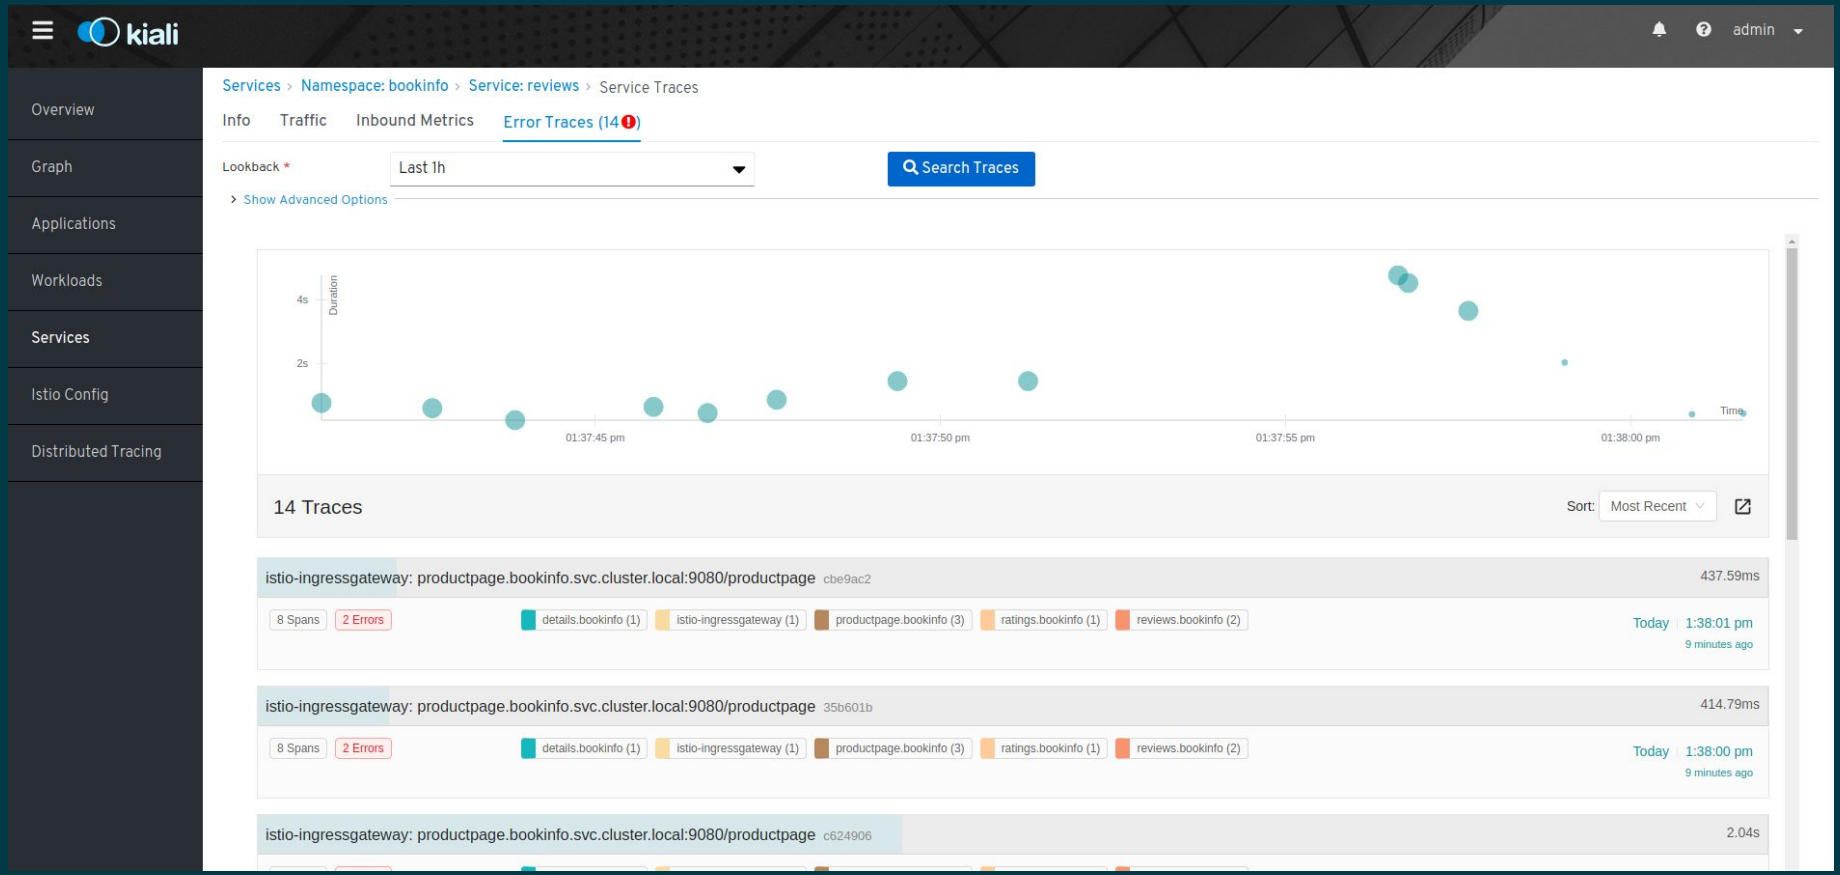

Visualizing security